

Probability and Statistics

Contents

1	Probability	3
1.1	Basic definitions	3
1.2	Main probability distributions	4
1.3	Waiting time distributions	8
1.4	Conditional probability	9
1.5	Product measures	11
1.6	Moments	11
1.7	Stochastic processes	13
1.8	Convergence	14
2	Statistics	16
2.1	Confidence regions	22
2.2	Statistical decision	23
2.3	Tests for Gaussian models	27
2.4	Goodness of fit and independence tests	30
2.5	Non-parametric methods	32
2.6	Epidemiology	36
2.7	Data reduction	39
2.8	Bayesian statistics	40
3	Regression analysis	41
3.1	Linear models	41
3.2	Analysis of variance (ANOVA)	45
4	Survival analysis	47
5	Design of experiments	49
6	Misc	50

The aim of these notes is to summarize a large number of concepts in statistics, both as a short reference and as a quick overview of the subject for people familiar with mathematical notation. It is simply a summary and reorganization from various classical books on these topics.

Chapter 1

Probability

1.1 Basic definitions

Let $\emptyset \neq \Omega$ be a set. A σ -algebra (on Ω) is a system $\mathcal{F} \subset \mathcal{P}(\Omega)$ with $\Omega \in \mathcal{F}$, closed under complements and countable unions. (Ω, \mathcal{F}) is an *event space*.

(We can't take the whole $\mathcal{P}(\Omega)$ because it's too big, Vitali's theorem: we can't assign a probability.)

To build our model we start from simple sets whose probabilities we can assign easily, then we generate a σ -algebra: a subset $\mathcal{G} \subset \mathcal{P}(\Omega)$ *generates* a unique σ -algebra $\sigma(\mathcal{G})$ on Ω such that $\mathcal{G} \subset \sigma(\mathcal{G})$.

The choice is often canonical:

1. When Ω is at most countable, we will choose $\mathcal{F} = \mathcal{P}(\Omega)$.
2. When $\Omega = \mathbb{R}^n$ we will use the *Borel σ -algebra*

$$\mathcal{B}_{\mathbb{R}^n} := \sigma \left(\left\{ \prod_{i=1}^n [a_i, b_i] \mid a_i < b_i \in \mathbb{Q} \right\} \right).$$

Every $A \in \mathcal{B}_{\mathbb{R}^n}$ is called a *Borel set*. If $\Omega \subset \mathbb{R}^n$, we set $\mathcal{B}_\Omega := \{A \cap \Omega \mid A \in \mathcal{B}_{\mathbb{R}^n}\}$.

3. Given $(E_i, \mathcal{E}_i)_{i \in I}$, we define the *product σ -algebra* on $\prod_{i \in I} E_i$ as

$$\bigotimes_{i \in I} \mathcal{E}_i := \sigma \left(\left\{ \pi_i^{-1}(A_i) \mid i \in I, A_i \in \mathcal{E}_i \right\} \right).$$

Probability measure (or distribution) P on an event space (Ω, \mathcal{F}) is a function $P: \mathcal{F} \rightarrow [0, 1]$ satisfying the normalization property $P(\Omega) = 1$ and σ -additivity for pairwise disjoint events. (Ω, \mathcal{F}, P) is a *probability space*.

Theorem 1.1.1. P satisfies the following: $P(\emptyset) = 0$, $P(A \cup B) = P(A) + P(B) - P(A \cap B)$, monotonicity, σ -subadditivity and σ -continuity $P(A_n) \rightarrow P(A)$ if $A_n \nearrow A$ or $A_n \searrow A$.

Theorem 1.1.2. Let $(\Omega, \mathcal{F} = \sigma(\mathcal{G}), P)$ where $\mathcal{G} \subset \mathcal{P}(\Omega)$ is an intersection-stable system. Then P is uniquely determined by its restriction $P|_{\mathcal{G}}$ to \mathcal{G} .

By Carathéodory's theorem, a function P defined on a suitable system \mathcal{G} may be extended to a probability measure on $\sigma(\mathcal{G})$.

Probability density (function) for P is a function $f: \Omega \rightarrow [0, \infty)$ such that, for all A in \mathcal{F} , if Ω is discrete (resp. Borel), one has

$$P(A) = \sum_{\omega \in A} f(\omega) \quad \left(\text{resp. } P(A) = \int_A f(x) dx \right).$$

It is called respectively a discrete and a Lebesgue density. A discrete density is also called probability mass function (PMF).

In the discrete case there is a 1:1 correspondence between the set of probability measures on $(\Omega, \mathcal{P}(\Omega))$ and the set of $f: \Omega \rightarrow [0, 1]$ such that $\sum_{\omega \in \Omega} f(\omega) = 1$. In the continuous case $\Omega \subset \mathbb{R}^n$, every measurable $f: \Omega \rightarrow [0, \infty)$ with $\int_{\Omega} f(x) dx = 1$ determines a unique probability measure on $(\Omega, \mathcal{B}_{\Omega})$, but not conversely (not every continuous measure has a density; also, two densities differing on a measure-zero set give the same probability).

Distribution of a random variable A *random variable* from (Ω, \mathcal{F}) to (Ω', \mathcal{F}') is a map $X: \Omega \rightarrow \Omega'$ which is measurable, i.e. satisfying $X^{-1}(A') \in \mathcal{F}$ for all $A' \in \mathcal{F}'$ (enough to check for a generator of \mathcal{F}'). A random variable X from (Ω, \mathcal{F}, P) to (Ω', \mathcal{F}') defines a probability measure P_X on (Ω', \mathcal{F}') by the rule

$$P_X(A') := P(X^{-1}(A')) = P(X \in A') \quad \text{for } A' \in \mathcal{F}',$$

called the *distribution of X under P* . Two random variables are *identically distributed* if they have the same distribution.

A probability measure also called distribution, indeed it's the distribution of the random variable id_{Ω} .

Random variables can be used to reduce observation depth.

Distribution density a density f of P_X . It links these notions together:

$$P(X \leq c) = F_X(c) = \int_{-\infty}^c f(x) dx.$$

(Cumulative) distribution function (CDF) for $(\mathbb{R}, \mathcal{B}, P)$ is the function

$$F_P: \mathbb{R} \rightarrow [0, 1], \quad F_P(c) = P((-\infty, c]).$$

For $X: \Omega \rightarrow \mathbb{R}$ random variable on (Ω, \mathcal{F}, P) , its CDF is defined as

$$F_X = F_{P \circ X^{-1}}: \mathbb{R} \rightarrow [0, 1], \quad F_X(c) = P(X \leq c).$$

Conversely, every increasing right-continuous $F: \mathbb{R} \rightarrow [0, 1]$ with $F(-\infty) = 0$ and $F(+\infty) = 1$ is F_X for $X: (0, 1) \rightarrow \mathbb{R}, u \mapsto \inf\{c \in \mathbb{R} \mid F(c) \geq u\}$. So, every probability measure P on \mathbb{R} is uniquely determined by F_P .

1.2 Main probability distributions

Bose-Einstein distribution n indistinguishable particles to be distributed in N distinguishable cells. $\Omega = \{(k_1, \dots, k_N) \in \mathbb{Z}_+^N: \sum_{j=1}^N k_j = n\}$ has cardinality $\binom{n+N-1}{n}$.

Continuous uniform distribution $\Omega \subset \mathbb{R}^n$ a Borel set with n -dimensional volume $\lambda^n(\Omega) := \int 1_\Omega(x) dx =: \int_\Omega dx$. The measure with constant density $f(x) = 1/\lambda^n(\Omega)$.

Urn models can be labeled/colored - ordered/unordered - with/without replacement. A useful trick is reducing the observation depth via random variables to compute the probability on the target space as the distribution of the random variable. N total balls, sample of size n .

Discrete uniform Replacement, ordered, labeled.

$$\Omega = \{1, \dots, N\}^n \ni (\omega_1, \dots, \omega_n) = \omega.$$

Discrete uniform distribution U_Ω has density $f(\omega) = 1/N^n$.

n -fold product, Bernoulli Replacement, ordered, colored.

If the set E of color classes $C_i \subset \{1, \dots, N\}$ has discrete density f , then the n -fold product density of f is the density on $\bar{\Omega} = E^n$ given by

$$f^{\otimes n}(C_1, \dots, C_n) = \prod_{i=1}^n f(C_i) = \prod_{i=1}^n \frac{|C_i|}{N}.$$

If $E = \{0, 1\}$ and $f(1) = p \in [0, 1]$ we call it the *Bernoulli distribution* for n trials, when and takes the form $p^k(1-p)^{n-k}$ where k is the number of 1's.

To see this we use a random variable

$$\begin{aligned} X: \{1, \dots, N\}^n &= \Omega \rightarrow \bar{\Omega} = E^n \\ \omega &\mapsto (C_{\omega_1}, \dots, C_{\omega_n}) \end{aligned}$$

and the probability \bar{P} on $\bar{\Omega}$ which we are interested in is given by the following P probability (the uniform probability on Ω):

$$P(X = (C_{i_1}, \dots, C_{i_n})) = P(C_{i_1} \times \dots \times C_{i_n}) = \frac{|C_{i_1}| \cdots |C_{i_n}|}{|\Omega|}.$$

Multinomial, binomial... Replacement, unordered, colored.

$\mathcal{C} = \{C_1, \dots, C_r\}$ set of color classes.

$\hat{\Omega} = \{\mathbf{k} = (k_1, \dots, k_r) \in \mathbb{Z}^{\mathcal{C}} \mid \sum k_i = n\}$, find its probability measure. We start from the n -fold and forget the ordering:

$$\begin{aligned} S: \mathcal{C}^n &= \bar{\Omega} \rightarrow \hat{\Omega} \\ \bar{\omega} &\mapsto (S_C(\bar{\omega}))_{C \in \mathcal{C}} = \left(\sum_{i=1}^n 1_{\{C\}}(\bar{\omega}_i) \right)_{C \in \mathcal{C}} \end{aligned}$$

is the histogram of the sample $\bar{\omega}$ and S_C counts the occurrences of color C .

$$\begin{aligned} \bar{P}(S = \mathbf{k}) &= \sum_{\bar{\omega}: S(\bar{\omega}) = \mathbf{k}} \prod_{i=1}^n \frac{|\bar{\omega}_i|}{N} = \frac{n!}{k_1! \cdots k_r!} \prod_{i=1}^r \left(\frac{|C_i|}{N} \right)^{k_i} = \\ &=: \binom{n}{\mathbf{k}} \prod_{i=1}^r f(C_i)^{k_i} =: \mathcal{M}_{n,f}(\{\mathbf{k}\}) \end{aligned}$$

is the *multinomial* distribution (note the multinomial coefficient above).

Special case: the binomial distribution. $\mathcal{C} = \{0, 1\}$, $\widehat{\Omega} = \{(k, n-k) \mid k \in \mathbb{Z}\}$ is identified with $\Omega' = \{0, \dots, n\}$. Let $f(1) =: p$, then the distribution becomes

$$\mathcal{B}_{n,p}(\{k\}) := \binom{n}{k} p^k (1-p)^{n-k}.$$

Replacement, unordered, labeled: it is a special case of the above, when $r = n$ and $k_i = 1$ for all i . We get $n!/N^n$.

No replacement, ordered, labeled. $\Omega_{\neq} = \{\omega \in \{1, \dots, N\}^n \mid \omega_i \neq \omega_j \text{ for } i \neq j\}$ simply has the uniform distribution, $f(\omega) = (N-n)!/N!$.

No replacement, unordered, labeled.

$$\widetilde{\Omega} = \{\omega \subset \{1, \dots, N\} \mid |\omega| = n\}$$

The distribution turns out to be the uniform one:

$$f(\omega) = n! \frac{(N-n)!}{N!} = \frac{1}{\binom{N}{n}} = \frac{1}{|\widetilde{\Omega}|}$$

(it is obvious, formally can be seen using the random variable $\Omega_{\neq} \rightarrow \widetilde{\Omega}$).

Hypergeometric No replacement, unordered, colored.

$\mathcal{C} = \{C_1, \dots, C_r\}$ with $|C_i| = N_i$, $\sum N_i = N$.

$$\left\{ \zeta \subset \{1, \dots, N\} \mid |\zeta| = n \right\} =: \widetilde{\Omega} \xrightarrow{T} \widehat{\Omega} := \left\{ \mathbf{k} = (k_1, \dots, k_r) \in \mathbb{Z}^{\mathcal{C}} \mid \sum k_i = n \right\}$$

$$\zeta \mapsto (|\zeta \cap C_i|)_{i=1}^r$$

Since $\zeta \mapsto (\zeta \cap C_i)_{i=1}^r$ is a bijection $\{T = \mathbf{k}\} \rightarrow \prod_{i=1}^r \{\zeta_i \subset C_i \mid |\zeta_i| = k_i\}$, we get

$$\mathcal{U}_{\widetilde{\Omega}}(T = \mathbf{k}) = \frac{\prod_i \binom{N_i}{k_i}}{\binom{N}{n}} =: \mathcal{H}_{n, \mathbf{N}}(\{\mathbf{k}\}),$$

the *hypergeometric* distribution of parameters n , $\mathbf{N} = (N_i)_{i=1}^r$.

Multinomial approximation of hypergeometric distributions Let E be finite and f be a discrete density on E . If $N \rightarrow \infty$, $N_a \rightarrow \infty$ and $N_a/N \rightarrow f(a)$ for all $a \in E$, then pointwise $\mathcal{H}_{n,N} \rightarrow \mathcal{M}_{n,f}$. Intuitively: as the sample space increases it becomes irrelevant whether there is replacement or not; easy proof using asymptotic approximation.

Poisson approximation of binomial distributions Let $p_n \in [0, 1]$ be a sequence with $np_n \rightarrow \lambda > 0$. Then for every $k \in \mathbb{N}$ we have

$$\mathcal{P}_{\lambda}(\{k\}) := \lim_{n \rightarrow \infty} \mathcal{B}_{n,p_n}(\{k\}) = e^{-\lambda} \lambda^k / k!$$

which we call the *Poisson distribution* on $\Omega = \mathbb{Z}_+$ (by writing the exponential series expansion one sees the limit sums to 1 and thus defines a discrete density on \mathbb{Z}_+).

Proven similarly as the multinomial approximation, plus using the usual exponential approximation formula.

The Poisson probability measure $\mathcal{P}_{\alpha t}$ measures the probability of k insurance claims in the interval $(0, t]$ ($p_n = \alpha t/n$, so $np_n \rightarrow \alpha t =: \lambda$) with α the average number of claims per unit of time.

Theorem 1.2.1. For $n \in \mathbb{N}$ and $p \in (0, 1)$ the following error bound holds,

$$\sum_{k \geq 0} |\mathcal{B}_{n,p}(\{k\}) - \mathcal{P}_{np}(\{k\})| \leq 2np^2$$

Normal distribution Let $v > 0$, $\Omega_N := B(0, \sqrt{vN}] \subset \mathbb{R}^N$, $P_N = U_{\Omega_N}$, and the projection $\pi_1 : \Omega_N \rightarrow \mathbb{R}$ onto the first coordinate. Then

$$\lim_{N \rightarrow \infty} P_N(a \leq X_1 \leq b) = \int_a^b \frac{e^{-x^2/2v}}{\sqrt{2\pi v}} dx$$

and we define the *normal (or Gauss) distribution* $N_{m,v}$ with mean $m \in \mathbb{R}$ and variance $v > 0$ as the probability measure on $(\mathbb{R}, \mathcal{B})$ with density (shown to be a density)

$$\varphi_{m,v}(x) = \frac{e^{-(x-m)^2/2v}}{\sqrt{2\pi v}}$$

The mean (resp. variance) of a sum is the sum of the means (resp. variances): if $X \sim \mathcal{N}(m_X, v_X)$ and $Y \sim \mathcal{N}(m_Y, v_Y)$, then $(X+Y) \sim \mathcal{N}(m_X+m_Y, v_X+v_Y)$.

The following areas will be used in statistics:

$\int_{-\sigma}^{\sigma} \varphi_{0,1} dx \approx 68\%$ (meaning that roughly 68% of the data are within 1 SD from the average)

$\int_{-2\sigma}^{2\sigma} \varphi_{0,1} dx \approx 95\%$

$\int_{-3\sigma}^{3\sigma} \varphi_{0,1} dx \approx 99.7\%$.

Exponential families An exponential family with parameters η is

$$p(\mathbf{x}|\eta) = g(\eta)h(\mathbf{x})\exp(\eta^T \mathbf{u}(\mathbf{x}))$$

where $\mathbf{u}(\mathbf{x})$ is a function of \mathbf{x} and $g(\eta)$ normalizes the distribution making it sum to 1. Examples include:

- Bernoulli

$$f(x, p) = p^x(1-p)^{1-x} = (1-p)\exp\left(\ln\left(\frac{p}{1-p}\right)x\right) = \sigma(-\eta)\exp(\eta x)$$

where $\sigma(\eta) = \frac{1}{1+\exp(-\eta)}$ is the logistic sigmoid function.

- Multinomial...

We have $-\nabla \ln g(\eta) = E(\mathbf{u}(\mathbf{x}))$ and the maximum likelihood estimator η_{ML} satisfies

$$-\nabla \ln g(\eta_{\text{ML}}) = \frac{1}{N} \sum_{n=1}^N \mathbf{u}(\mathbf{x}_n) \rightarrow E(\mathbf{u}(\mathbf{x}))$$

so $\sum_{n=1}^N \mathbf{u}(\mathbf{x}_n)$ is a sufficient statistic of the distribution.

Mixture distributions If p_i are distribution densities (or cumulative distribution functions), then the sum $\sum_i w_i p_i$ with weights summing to 1 also represents a distribution. The uncountable case is written as $f(x) = \int_{\Theta} p(x|\theta)q(\theta)$ where $p(x|\theta)$ are a family of distributions parametrized by θ , and θ is itself a random variable with distribution $q(\theta)$.

1.3 Waiting time distributions

Negative binomial (Pascal), geometric The waiting time until the n th success, i.e. number of failures before the n th success is

$$T_n: \{0, 1\}^{\mathbb{N}} \rightarrow \Omega = \mathbb{N}, \quad T_n(\omega) = \min\{k \in \mathbb{N} \mid \sum_{i=1}^k \omega_i = n\} - n$$

We want the probability of the n th success on the $(n+k)$ th draw; we can compute it as the distribution of T_n or directly as

$$\binom{n+k-1}{k} p^n (1-p)^k =: \binom{-n}{k} p^n (p-1)^k =: \overline{B}_{n,p}(\{k\})$$

where p is the probability of success and the binomial coefficient is the number of choices for the k failure times (the -1 comes from the last draw having to be a success). This is called the *negative binomial distribution* or the *Pascal distribution*. When $n = 1$ we call it the *geometric distribution*.

Gamma, exponential For a fixed α , the probability of at least k claims in the interval $(0, t]$ as a function of t is given by

$$\Gamma_{\alpha,k}((0, t]) = 1 - P_{\alpha,t}(\{0, \dots, k-1\}) = \int_0^t \frac{\alpha^k}{(k-1)!} s^{k-1} e^{-\alpha s} ds =: \int_0^t \gamma_{\alpha,k}(s) ds$$

hence we get a probability measure on $(0, \infty)$ with Lebesgue density $\gamma_{\alpha,k}$ (use Euler's Γ function

$$\Gamma(r) = \int_0^{\infty} s^{r-1} e^{-s} ds, \quad r > 0$$

to check that it is a density). We call this the *Gamma distribution* $\Gamma_{\alpha,k}$. It can be generalized to $\Gamma_{\alpha,r}$ for every real $r > 0$, replacing $(k-1)!$ with $\Gamma(r)$. If $r = 1$ we call it the *exponential distribution* (the first point i.e. "claim" is exponentially distributed).

Beta β_{n-r+1} is the distribution of the r th order statistic of an n -sample from a continuous uniform distribution. If we are to be given n fixed points in the unit interval $(0, 1)$, i.e. some $\omega := (\omega_1, \dots, \omega_n) \in \Omega := (0, 1)^n$ with the uniform distribution, where is the r th point in order? (E.g. if we are expecting n deliveries, how long does it take to receive the r th?) The projection $\pi_i(\omega)$ is the instant of the i -th point. We see π_i as random variables on Ω . Define the random variables $\pi_{r:n}$ so that

$$\pi_{1:n} < \pi_{2:n} < \dots < \pi_{n:n} \quad \{\pi_{1:n}, \dots, \pi_{n:n}\} = \{\pi_1, \dots, \pi_n\}.$$

The *Beta distribution* is defined as

$$P(\pi_{r:n} \leq c) = \int_{-\infty}^c \beta_{r,n-r+1}(s) ds$$

where

$$\beta_{a,b}(s) = \frac{s^{a-1}(1-s)^{b-1}}{B(a,b)} \quad 0 < s < 1$$

where $B(a,b)$ is Euler's Beta function. In particular the distribution density of the first point is $\beta_{1,n}(s) = n(1-s)^{n-1}$ and the probability measure given by the density $\beta_{1,1}$ is the uniform probability on $[0,1]$.

1.4 Conditional probability

$P(B|A) := \frac{P(B \cap A)}{P(A)}$; $P(-|A)$ is the unique probability measure such that A is certain and it's proportional to P on subevents of A .

It is not related to causality, it can be interpreted as our evaluation of the situation after we have been *informed* that A has happened (but not after it *has* happened, in fact B could happen before A ! For example: draw a ball without looking, then draw another one and see its color, then we can re-evaluate the probability of the color of the first one.

Let $\Omega = \bigcup_{i \in I} B_i$ be a countable partition.

- Case-distinction formula: $P(A) = \sum_{i \in I} P(B_i)P(A|B_i)$.
- Bayes' formula: $P(B|A) = \frac{P(B)P(A|B)}{P(A)}$ or, more in general,

$$P(B_k|A) = \frac{P(B_k)P(A|B_k)}{\sum_{i \in I} P(B_i)P(A|B_i)}.$$

- Multiplication rule: given events A_1, \dots, A_n , we have

$$P(A_1 \cap \dots \cap A_n) = P(A_1)P(A_2|A_1) \dots P(A_n|A_1 \cap \dots \cap A_{n-1}).$$

This helps us study the case when the distribution of a first random variable X_1 is known and X_k is known when X_1, \dots, X_{k-1} are known. Along with the next results it is useful in the construction of multi-stage models.

Evaluation of medical tests Ω = population, $P = U_\Omega$, $D \subset \Omega$ people with the disease, $P \subset \Omega$ people testing positive.

$P(D)$ = prevalence of the disease, $P(P|D)$ = test sensitivity, $P(P|D^c) = (1 - \text{specificity})$.

Positive correctness: $P(D|P) = \frac{P(D)P(P|D)}{P(D)P(P|D) + P(D^c)P(P|D^c)}$ can be low if the prevalence is low, even though the sensitivity is high.

Negative correctness: $P(D^c|P^c) = \dots$

Three doors problem Probability is $2/3$ if we assume the host will not open the chosen door, or $1/2$ otherwise.

Theorem 1.4.1 (Construction of probability measures via conditional probabilities). Let $\Omega_1, \dots, \Omega_n$ be countable, f_1 a density on Ω_1 , $f_k|_{\omega_1, \dots, \omega_{k-1}}$ a density on Ω_k for $k \geq 2$ and $\omega_i \in \Omega_i$ for $i = 1, \dots, k-1$. Let $\Omega = \prod_{i=1}^n \Omega_i$, π_i be the i -th projection.

Then, there is a unique probability measure P on $(\Omega, \mathcal{P}(\Omega))$ with the properties:

1. $P(\pi_1 = \omega_1) = f_1(\omega_1)$ for all $\omega_1 \in \Omega_1$,
2. $P(\pi_k = \omega_k \mid \pi_1 = \omega_1, \dots, \pi_{k-1} = \omega_{k-1}) = f_{k|\omega_1, \dots, \omega_{k-1}}(\omega_k)$.

This P is given by

$$P(\{(\omega_1, \dots, \omega_n)\}) = f_1(\omega_1) f_{2|\omega_1}(\omega_2) \cdots f_{n|\omega_1, \dots, \omega_{n-1}}(\omega_n).$$

In the case of an infinite product space $\Omega = \prod_{i \geq 1} \Omega_i$, replace $\mathcal{P}(\Omega)$ by $\bigotimes_{i \geq 1} \mathcal{P}(\Omega_i)$. Then the theorem becomes: there exists a unique probability measure P such that, for all $k \geq 1$ and $\omega_i \in \Omega_i$,

$$P(\pi_1 = \omega_1, \dots, \pi_k = \omega_k) = f_1(\omega_1) f_{2|\omega_1}(\omega_2) \cdots f_{k|\omega_1, \dots, \omega_{k-1}}(\omega_k).$$

The finite case: the stated expression is identical to the multiplication rule above, under the conditions of the theorem, proving uniqueness. Then we check that it satisfies the two properties.

Independence

- A family $(A_i)_{i \in I}$ of events is called independent if, for every finite subset $\emptyset \neq J \subset I$, we have

$$P\left(\bigcap_{i \in J} A_i\right) = \prod_{i \in J} P(A_i).$$

- Since we are also interested in independence of subexperiments in addition to events, we define a family $(X_i)_{i \in I}$ of random variables $\Omega \rightarrow \Omega_i$ independent if the family $(\{X_i \in B_i\})_{i \in I}$ is independent for every choice $B_i \in \mathcal{F}_i$.

Note this is stronger than pairwise independence. Toss two coins and consider events: first toss is heads, second toss is heads, both tosses give same result. They are pairwise independent but not independent as a family.

Independence despite causality: e.g., rolling two dice, A = sum of points is 7, B = first dice shows 6 are independent although the sum depends causally on the value of the first dice.

As a consequence of the uniqueness theorem, to check independence of random variables it is enough to choose $B_i \in \mathcal{G}_i$ where \mathcal{G}_i is an intersection-stable generator of \mathcal{F}_i .

A family of events is independent iff the corresponding family of indicator functions is independent.

Criterion for independence of finitely many random variables. Consider a finite sequence of random variables $Y_i: \Omega \rightarrow \Omega_i$ ($i = 1, \dots, n$) on a probability space (Ω, \mathcal{F}, P) . In the discrete case of countable Ω_i 's, $(Y_i)_i$ is independent iff $P(Y_1 = \omega_1, \dots, Y_n = \omega_n) = \prod_{i=1}^n P(Y_i = \omega_i)$ for all $\omega_i \in \Omega_i$. In the real case $\Omega_i = \mathbb{R}$, the family is independent iff $P(Y_1 \leq c_1, \dots, Y_n \leq c_n) = \prod_{i=1}^n P(Y_i \leq c_i)$ for all $c_i \in \mathbb{R}$. (The "only if" direction is trivial in both cases).

For example, the n -fold product measure $f^{\otimes n}$ on E^n makes the projection random variables independent.

1.5 Product measures

We want to construct a model for infinite coin tosses, so we need to prove existence of infinitely many independent random variables and work with infinite product σ -algebras.

Theorem 1.5.1 (Construction of independent random variables with prescribed distributions). Let $(\Omega_i, \mathcal{F}_i, P_i)_{i \in I}$ be probability spaces, I countable. Then there exists (Ω, \mathcal{F}, P) and independent $Y_i: \Omega \rightarrow \Omega_i$ such that $P \circ Y_i^{-1} = P_i$ for all $i \in I$.

Corollary 1.5.2 (Existence of infinite product measures). Let $(\Omega_i, \mathcal{F}_i, P_i)_{i \in I}$ be probability spaces, I countable. Then there exists a unique probability measure P on $(\Omega := \prod_i \Omega_i, \mathcal{F} := \bigotimes_i \mathcal{F}_i)$ such that for all finite $J \subset I$ and all $A_i \in \mathcal{F}_i$ one has

$$P(\pi_i \in A_i \text{ for all } i \in J) = \prod_{i \in J} P_i(A_i)$$

i.e. the projections $\pi_i: \Omega \rightarrow \Omega_i$ are independent with distribution P_i . We denote this P by $\bigotimes_i P_i$ and call it the *product measure* of the P_i 's.

The proof of the above corollary also shows that independence and product measures are closely related: a countable family Y_i of random variables is independent iff $P \circ (Y_i)_{i \in I}^{-1} = \bigotimes_{i \in I} P \circ Y_i^{-1}$ (the joint distribution equals the product of marginal distributions).

The density of a finite product measure is the product of the single densities, if they exist. An infinite product measure never admits a density, since there is no Lebesgue measure on $\mathbb{R}^{\mathbb{N}}$.

Convolution Measures Q_1, Q_2 on \mathbb{R} , then the distribution of the sum $Q_1 + Q_2$ is the convolution $Q_1 * Q_2 := (Q_1 \otimes Q_2) \circ A^{-1}$ where $A: \mathbb{R}^2 \rightarrow \mathbb{R}$ sends $(x_1 + x_2) \mapsto x_1 + x_2$.

1.6 Moments

Expectation of a real random variable is defined as the first moment (when it exists)

$$\mu := E(X) := \int_{\Omega} X dP = \int_{\mathbb{R}} x dP_X$$

where the equality follows from the change of variable theorem

$$\int_{\Omega} (g \circ X) dP = \int_{\mathbb{R}} g dP_X$$

(this is the law of the unconscious statistician, LOTUS, by which we can write the expectation of $E(gX)$ without the distribution of gX). In the Lebesgue case we can further write $dP_X = dF_X = f_X dx$.

If $E|X| < \infty$ we say X has an expectation and write $X \in \mathcal{L}^1$.

If P on $(\mathbb{R}, \mathcal{B})$ then we define $E(P) = E_P(\text{Id}_{\mathbb{R}})$ and $E(P) = \text{Var}_P(\text{Id}_{\mathbb{R}})$ if they exist.

E is linear and commutes with products for independent variables.

Variance If $X \in \mathcal{L}^2$ we define its variance as the second central moment

$$\text{Var}(X) := \sigma^2 := \mathbb{E}(X - \mathbb{E}X)^2 = \mathbb{E}(X^2) - \mathbb{E}(X)^2,$$

and σ is called standard deviation. We define the covariance of $X, Y \in \mathcal{L}^2$ as

$$\text{Cov}(X, Y) := \mathbb{E}((X - \mathbb{E}X)(Y - \mathbb{E}Y)) = \mathbb{E}(XY) - \mathbb{E}(X)\mathbb{E}(Y).$$

We have $\text{Cov}(aX + b, cY + d) = ac\text{Cov}(X, Y)$, in particular $\text{Var}(aX + b) = a^2\text{Var}(X)$, and

$$\text{Var}\left(\sum X_i\right) = \sum \text{Var}(X_i) + \sum_{i \neq j} \text{Cov}(X_i, X_j)$$

(for X_i pairwise uncorrelated, i.e. when there are no covariances, this is called Bienaymé's identity).

The correlation is defined as $f(X, Y) = \text{Cov}(X, Y)/\sigma_X\sigma_Y \in [-1, 1]$; it is a measure of association (it is the slope of linear regression line; cfr. SD line, axis of symmetry of the scatter plot). Independent implies uncorrelated.

Law of total expectation and variance $\mathbb{E}(X) = \mathbb{E}(\mathbb{E}(X|Y))$ is the law of total expectation. This can be used to prove the law of total variance $\text{Var}(X) = \mathbb{E}(\text{Var}(X|Y)) + \text{Var}(\mathbb{E}(X|Y))$.

General moments The n -th moment of X is $\mathbb{E}X^n$; the n -th central moment is $\mu_n := \mathbb{E}(X - \mu)^n$; the n -th standardized (or normalized) moment is $\frac{\mathbb{E}(X - \mu)^n}{\sigma^n}$. The moment generating function (mgf) is $M_X(t) = \mathbb{E}\exp(tX)$ and then the n th moment is its n th derivative evaluated at 0. If X and Y are independent, then $M_{X+Y}(t) = M_X(t)M_Y(t)$ whence we get, for example, that the mean and variance of a sum of normal distributions are the sum of the means and sum of the variances. The sequence of moments does not determine a distribution uniquely in general, unless the random variables have bounded support (but equality of the first few moments can be used to approximate a distribution...).

The following are examples of higher-order statistics (using the third or higher power of a sample).

Skewness is the third standardized moment $\frac{\mathbb{E}(X - \mu)^3}{\sigma^3}$, which measures the symmetry of the distribution: negative (or left) skew means the left tail is longer and the mass of the distribution is concentrated to the right, and symmetrically for positive (or right) skew.

Skewness is not directly related to the relationship between mean and median (though often the mean is to the left of the median in negative skew), unlike the old definition of non-parametric skew: $(\mu - \nu)/\sigma$ where ν is the median. Another possible definition is Pearson's mode skewness $(\mu - \text{mode})/\sigma$.

Kurtosis is the fourth standardized moment $K(X) := \frac{\mathbb{E}(X - \mu)^4}{\sigma^4}$; excess kurtosis is $K - 3$ since $K = 3$ for a normal distribution. A distribution with positive excess kurtosis is called leptokurtic (fat-tailed, has heavy tails, more mass on the tails of its support than a normal distribution, more extreme outliers, otherwise platykurtic (flatter, e.g. uniform).

Mode A mode of a distribution is any local maximum of the density function (or probability mass function for discrete distributions). The distribution is unimodal if the maximum is unique, multimodal otherwise. For symmetric unimodal distributions, the mean (if defined), median and mode coincide.

1.7 Stochastic processes

A stochastic process on a probability space (Ω, \mathcal{F}, P) is a collection

$$X = \{X_t: (\Omega, \mathcal{F}) \rightarrow (S, \Sigma) \mid t \in T\}.$$

A filtration of \mathcal{F} is a collection $(\mathcal{F}_t)_t$ of sub- σ -algebras such that $\mathcal{F}_s \subset \mathcal{F}_t$ if $s \leq t$ (the index set T is totally ordered). Standard example:

$$\mathcal{F}_t = \sigma(X_s \mid s \leq t) = \sigma(X_s^{-1}(A) \mid A \in \Sigma, s \leq t)$$

called the natural filtration of \mathcal{F} wrt X : it records past behavior of X .

X is adapted (or non-anticipating) to a filtration $(\mathcal{F}_t)_t$ if it cannot see into the future: $X_t: \Omega \rightarrow S$ is (\mathcal{F}_t, Σ) -measurable for each t .

A stopping time wrt $(\mathcal{F}_t)_t$ is a random variable $\tau: \Omega \rightarrow T \cup \{+\infty\}$ such that $\{\tau \leq t\} \in \mathcal{F}_t$ for all t .

A process X adapted to $(\mathcal{F}_t)_t$ has the Markov property if, for all $A \in \Sigma$ and $s < t$, one has

$$P(X_t \in A \mid \mathcal{F}_s) = P(X_t \in A \mid X_s),$$

alternatively $E(f(X_t) \mid \mathcal{F}_s) = E(f(X_t) \mid \sigma(X_s))$

Intuitively, the future only depends on the present state and not the past.

Memorylessness for a random variable is the property $P(X > s + t \mid X > s) = P(X > t)$ for all s, t integer or real numbers (depending on whether X is discrete or continuous). The only such distributions are the geometric (discrete) and exponential (continuous): defining the survival function $S(t) = P(X > t)$, the property can be written as $S(t + s) = S(t)S(s)$ whence $S(q) = S(1)^q$ for rational q , and the only continuous S satisfying this is $S(x) = \exp(-\ln S(1)x)$. (See survival analysis). It's weaker than the Markov property.

Martingale A stochastic process Y_i is a martingale wrt another one X_i if Y_i has finite expectation and the conditional expected value of the next observation given the previous observations, is equal to the latest observation:

$$\begin{aligned} E(Y_{n+1} \mid X_1, \dots, X_n) &= Y_n \text{ (discrete)} \\ E(Y_t \mid \{X_{s'}, s' \leq s\}) &= Y_s \text{ for all } s \leq t \text{ (continuous)} \\ E(Y_t \mid \mathcal{F}_s) &= Y_s \text{ for all } s \leq t \text{ (continuous)} \end{aligned}$$

In the case $X = Y$, we just say X is a martingale. Examples: unbiased random walk, gambler's capital in fair games.

1.8 Convergence

- Convergence in distribution: $F_n(x) \rightarrow F(x)$ for all x where $F(x)$ is continuous (the CDFs of X_n, X resp.), denoted by $X_n \rightsquigarrow X$. If F is continuous this is equivalent to uniform convergence $\|F_n - F\|_\infty \rightarrow 0$.
- Convergence in probability: for every ε , $P(|X_n - X| > \varepsilon) \rightarrow 0$, denoted by $X_n \xrightarrow{P} X$.
- Convergence in \mathcal{L}^p : $E|X_n - X|^p \rightarrow 0$, denoted $X_n \xrightarrow{\mathcal{L}^p} X$.
- Almost sure convergence: $P(\{\omega \in \Omega \mid X_n(\omega) \rightarrow X(\omega)\}) = 1$ (i.e. point-wise almost-everywhere), denoted by $X_n \xrightarrow{\text{a.s.}} X$.

Either of $\xrightarrow{\text{a.s.}}$ or $\xrightarrow{\mathcal{L}^p}$ implies \xrightarrow{P} which in turn implies \rightsquigarrow .

A random variable is asymptotically normal if it converges in distribution to a normally distributed random variable.

Weak law of large numbers (WLLL)

- \mathcal{L}^2 version: if $X_i \in \mathcal{L}^2$ are pairwise uncorrelated with variance bounded by v , then $\frac{1}{n} \sum_{i=1}^n (X_i - EX_i) \xrightarrow{P} 0$; more precisely, for all $\varepsilon > 0$,

$$P\left(\left|\frac{1}{n} \sum_{i=1}^n (X_i - EX_i)\right| \geq \varepsilon\right) \leq \frac{v}{n\varepsilon^2}$$

(sample mean converges in probability to the average of the expectations).
If they are iid, we may use the following version:

- \mathcal{L}^1 version: if $X_i \in \mathcal{L}^1$ pairwise iid with $E(X_i) =: \mu$, then we can write the result as $\frac{1}{n} \sum_{i=1}^n X_i \xrightarrow{P} \mu$ (sample mean converges in probability to the mean).

\mathcal{L}^2 proof: the argument inside the absolute value is \mathcal{L}^2 with mean 0 and variance $\frac{1}{n^2} \sum \text{Var}(X_i) \leq v/n$, so the theorem follows immediately from Chebyshev's inequality: $P(|Y - EY| \geq \varepsilon) \leq \text{Var}(Y)/\varepsilon^2$ for all $Y \in \mathcal{L}^2$ and $\varepsilon > 0$.

Note: Chebyshev's inequality implies that the proportion of samples within k standard deviations is $\geq 1 - 1/k^2$.

Strong law of large numbers (SLLL) Under the \mathcal{L}^2 hypotheses, convergence of $\frac{1}{n} \sum_{i=1}^n (X_i - EX_i) \rightarrow 0$ is almost sure (which implies convergence in probability too).

Central limit theorem $X_i \in \mathcal{L}^2$ iid with $E(X_i) = \mu$ and $\text{Var}(X_i) = \sigma^2 > 0$, then

$$\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \rightsquigarrow \mathcal{N}_{0,1}$$

that is, the sample mean of iid's tends to a normal distribution with mean and variance of the sample mean, regardless of the particular distribution.

Applying this to X_i Bernoulli we get as a special case the normal approximation of \mathcal{B} by De Moivre-Laplace, also known as local limit theorem. In this case we have convergence not only in distribution but even uniformly.

Normal approximation of \mathcal{B} Let $p \in (0, 1)$, $q = 1 - p$, $c > 0$, $x_n(k) = \frac{k - np}{\sqrt{npq}}$.

$$\max_{k: |x_n(k)| \leq c} \left| \frac{\mathcal{B}_{n,p}(\{k\})}{\varphi(x_n(k))/\sqrt{npq}} - 1 \right| \xrightarrow{n \rightarrow \infty} 0$$

This says that the histogram tends to a Gaussian curve; the integral form of this theorem states the same in terms of the area and is useful to avoid binomial computations: for $0 \leq k \leq l \leq n$ one has

$$\mathcal{B}_{n,p}(\{k, \dots, l\}) \xrightarrow{n \rightarrow \infty} \Phi(x_n(l + 1/2)) - \Phi(x_n(k - 1/2))$$

Chapter 2

Statistics

Statistical model A collection of probability spaces $(X, \mathcal{F}, P_\theta : \theta \in \Theta)$, where the P_θ are probability measures of the same class. It is *parametric* if $\Theta \subset \mathbb{R}^d$, *discrete* if X is discrete, *continuous* if $X \in \mathcal{B}_{\mathbb{R}^n}$, $\mathcal{F} = \mathcal{B}_X$. A *standard* model is a model that is either discrete or continuous.

Statistics studies the following problem: by looking at some observations (represented by elements of X), how can we determine the right θ , i.e. the probability measure P_θ that controls our phenomenon?

Estimators A *statistic* is a random variable on (X, \mathcal{F}) .

A (*point*) *estimator* of θ is a statistic $\hat{\theta}: X \rightarrow \Theta$ (Θ must have a σ -algebra); its values are the *estimates*. Clearly, for this to be useful, we must choose a $\hat{\theta}$ that "resembles" θ in some way.

Maximum Likelihood Estimation The likelihood function for an outcome $x \in X$ is defined as $\mathcal{L}(-|x) := f(x| -)$, describing the probability of observing x under the various probabilities P_θ . Then we can estimate θ as the $\hat{\theta}(x)$ for which x has the largest $P_{\hat{\theta}(x)}$ -probability:

$$f(x|\hat{\theta}(x)) = \max_{\theta \in \Theta} f(x|\theta) \quad \text{i.e.} \quad \hat{\theta}(x) := \arg \max_{\theta \in \Theta} \mathcal{L}(-|x)$$

and we call $\hat{\theta}$, defined pointwise, a *maximum likelihood estimator (MLE)* for θ .

A likelihood ratio $f(X = x|\theta_1)/f(X = x|\theta_2)$ describes whether θ_1 or θ_2 is more likely.

Note on Bayesian statistics In the frequentist interpretation there is no probability over the parameter space Θ , whereas in Bayesian statistics, the prior and the posterior are probabilities on that space. We start by choosing a density $f(\theta)$ over $(\Theta, \mathcal{B}_\Theta^d)$ before any data is seen, expressing our belief for each θ based on some assumptions; this is called prior probability. Then, after x is observed, we can update our belief by defining the posterior probability $f(\theta|x)$, using Bayes' theorem to take into account both the prior and the likelihood of x under θ :

$$\text{posterior} = f(\theta|x) := \frac{f(x|\theta)f(\theta)}{f(x)} = \frac{\text{likelihood} \times \text{prior}}{\text{evidence}}$$

where

$$f(x) := \int f(x|\theta)f(\theta)d\theta$$

is called the (model) evidence or marginal likelihood and acts as normalizing constant.

Maximum a-posteriori (MAP) estimation consists in maximizing $f(x|\theta)f(\theta)$ over θ , hence accounting for the prior distribution, instead of maximizing only the likelihood $f(x|\theta)$ as in MLE.

The evidence is often intractable (it requires integrating over all possible combinations of parameters); finding approaches to solve this problem is the objective of Bayesian statistics.

Sample mean and variance are the Gaussian MLEs The MLE for the n -fold Gaussian model $(\mathbb{R}^n, \mathcal{B}^n, \mathcal{N}_{\mu,v}^{\otimes n} : \mu \in \mathbb{R}, v > 0)$ is $(\bar{X}, V) : \mathbb{R}^n \rightarrow \mathbb{R} \times (0, \infty)$ where

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i \quad V = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

are called the *sample mean* and *sample variance* respectively.

Proof: must maximize $f(\mu, v, x) = (2\pi v)^{-n/2} \exp(-\sum_{i=1}^n \frac{(x_i - \mu)^2}{2v})$ wrt (μ, v) : we find \bar{X} by minimizing the mean squared error inside exp, which is obtained for the sample mean by Pythagoras' theorem

$$\frac{1}{n} \sum (x_i - \mu)^2 = \frac{1}{n} \sum (x_i - \bar{X})^2 + (\bar{X} - \mu)^2;$$

then we find V by differentiating $\log f_x(\bar{X}, v)$ wrt v ...

Bias What are some desirable requirements for an estimator? An estimator T (or a sequence of estimators T_n depending on observations X_1, \dots, X_n) is *unbiased* if

$$\text{Bias} := E_\theta(T) - \theta = 0 \quad \text{for all } \theta \in \Theta.$$

$\{T_n\}$ is *asymptotically unbiased* if $E_\theta(T_n) \rightarrow \theta$ as $n \rightarrow \infty$, for all $\theta \in \Theta$. For example, if T_n is such that $E_\theta(T_n) = \frac{n}{n+1}\theta$ then it's biased but asymptotically unbiased.

In a series of measurements, the bias is an error affecting all measurements in the same way, unlike chance errors. Unbiasedness assures that the values are typically centered around θ , but for this to be useful we should also know that they don't fluctuate too much, that is, that T has small variance. Note that a small variance is not useful in the case of a very biased estimator (because then the values would be centered around a wrong value).

We will now see that an MLE is not necessarily unbiased.

Unbiased estimators for mean and variance In a general n -fold product model $(\mathbb{R}^n, \mathcal{B}^n, P_\theta^{\otimes n} : \theta \in \Theta)$ with expectation and variance defined for all θ , the sample mean \bar{X} is an unbiased estimator of $\mu(\theta) := E(P_\theta)$ (immediate by

linearity). Reordering Pythagoras' theorem above and applying E, we compute the expectation of the sample variance

$$\begin{aligned} E(V) &= \frac{1}{n} \sum E((X_i - \mu)^2) - E((\bar{X} - \mu)^2) = \\ &= \frac{1}{n} \sum (\text{Var}(X_i - \mu) - E((X_i - \mu)^2)) - \text{Var}(\bar{X} - \mu) = \\ &= \frac{1}{n} \sum \text{Var}(X_i) - \frac{\text{Var}(X_i)}{n} = \frac{n-1}{n} \sigma^2 : \end{aligned}$$

we see it is underestimating the variance (because of the distance $(\bar{X} - \mu)^2$ in the Pythagoras expression); to get an unbiased estimator of $v(\theta) := \text{Var}(P_\theta)$, we must take the Bessel-corrected sample variance

$$V^* := S^2 := \frac{n}{n-1} V = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2.$$

Properties of M, V Distributions... (thm 5.3.1 Casella Berger page 248); sum of normal distributions is normal with... \bar{X} and V^* are independent random variables. $\bar{X} \sim \mathcal{N}(m, v/n)$ (using distribution of sum of normal variables), $(n-1)V^*/v \sim \chi_{n-1}^2$.

Pooled variance Given k groups of samples of cardinality n_1, \dots, n_k with pairwise different sample means m_1, \dots, m_k and same variance, the pooled sample variance is the unbiased estimator of the common variance defined as:

$$\begin{aligned} s_p^2 &= \sum_{i=1}^k \frac{(n_i - 1)s_i^2}{\sum_{i=1}^k (n_i - 1)} = \sum_{i=1}^k \frac{(n_i - 1)s_i^2}{n_1 + \dots + n_k - k} \\ &= \frac{1}{n - k} \sum_{i=1}^k \sum_{x_j \in G_i} (x_j - m_i)^2 \end{aligned}$$

($n - k$ degrees of freedom because of the k different sample means). It's a weighted average of the sample variances and it coincides with the simple average when all groups have the same cardinality.

$$s_p^2 / \sigma^2 \sim \chi_{(\sum n_i) - m}^2.$$

(maybe the left side should be multiplied by $n_1 + \dots + n_k - k$?)

- Do not confuse this with the combined variance:

$$\frac{n_1 s_1^2 + n_2 s_2^2 + n_1 (\bar{X}_1 - \bar{X})^2 + n_2 (\bar{X}_2 - \bar{X})^2}{n_1 + n_2}$$

where \bar{X}_i is the mean for group i and \bar{X} is the overall mean.

Behrens-Fisher problem It's the problem of testing the difference between the means of two independent samples *without* the assumption of equality of variances.

In general, given an \mathbb{R} -linear combination $\chi' := \sum_1^n k_i s_i^2$ of independent sample variances s_i with ν_i DoF (usually $\nu_i = n_i - 1$ and $k_i = \frac{1}{n_i - 1}$), this time with no assumption of equality of the σ_i 's, the distribution of χ' can't be expressed in general, but can be approximated by a χ distribution with DoF given by the Welch-Satterthwaite equation:

$$\nu_{\chi'} = \left(\sum_1^n k_i s_i^2 \right)^2 / \sum_1^n \frac{(k_i s_i^2)^2}{\nu_i}$$

Decision theory How to choose between different estimators, or measure the quality of an estimator? A loss function (or error function) $L: \Theta \times \Theta \rightarrow \mathbb{R}$ measures the distance $L(\theta, \hat{\theta}(x))$ e.g. squared error, absolute error, etc... The risk function (or expected prediction error) is defined as

$$R(\theta, \hat{\theta}) = E_{\theta}(L(\theta, \hat{\theta})) = \int L(\theta, \hat{\theta}(x)) f(x, \theta) dx$$

(E_{θ} is the expectation under the probability P_{θ}). This is the mean squared error (MSE) when the loss function is the square; in this case we can derive the expression for the bias-variance tradeoff:

$$\begin{aligned} \text{MSE}(T) &:= E_{\theta}((T - \theta)^2) = \text{Var}_{\theta}(T) + (E_{\theta}(T) - \theta)^2 = \\ &= \text{Var}_{\theta}(T) + \text{Bias}_T(\theta)^2 \end{aligned}$$

and this is the quantity that we aim to minimize, even at the cost of admitting some bias. For example, if we know that θ belongs to some range where the error of T' is smaller than the error of T , then we use T' even if its bias is bigger than that of T . (We will always suppose that our estimators belong to $\bigcap_{\theta \in \Theta} \mathcal{L}^2(P_{\theta})$ so that the variances are defined.)

An unbiased estimator T is called a *best estimator* if $\text{Var}_{\theta}(T) \leq \text{Var}_{\theta}(S)$ for all other unbiased estimators S and all $\theta \in \Theta$.

We may choose $\hat{\theta}$ so as to minimize either the maximum risk $\sup_{\theta} R(\theta, \hat{\theta})$ (a minimax rule) or the Bayes risk $\int R(\theta, \hat{\theta}) f(\theta) d\theta$ (a Bayes rule, $f(\theta)$ is a prior distribution).

In the regression setting we further average over " θ ", we write:

$$E(L) = \iint L(y, \hat{y}(\mathbf{x})) f(\mathbf{x}, y) d\mathbf{x} dy;$$

when L is the square function, this is minimized by the conditional mean $\hat{y}(\mathbf{x}) := \int y f(y|\mathbf{x}) dy = E_y[y|\mathbf{x}]$, the conditional expectation (chapter 3).

In the classification setting, we should minimize

$$E(L) = \int \sum_k L(C_k, \hat{C}_k(\mathbf{x})) f(\mathbf{x}, C_k) d\mathbf{x}.$$

If L is the 0-1 loss ("accuracy"), the minimum is given pointwise by the *Bayes classifier*

$$\hat{C}(\mathbf{x}) = \text{argmax}_{i=1}^k P(C_i | X = \mathbf{x}),$$

which we can estimate e.g. by K-nearest neighbors, Naive Bayes Classifier...

Method of moments Another way to estimate θ is to define $\hat{\theta}_n$ such that each moment $\int x^j dF_{\hat{\theta}_n}$ equals the sample moment $\frac{1}{n} \sum_{i=1}^n X_i^j$. These estimators are often biased.

Standard error The distribution of $\hat{\theta}_n$ is called the sampling distribution; its standard deviation is called the standard error $\text{SE}(\hat{\theta}_n) = \sqrt{\text{Var}(\hat{\theta}_n)}$. If we estimate it, we use notation $\widehat{\text{SE}}(\hat{\theta}_n)$. We estimate $\text{Var}_F(\hat{\theta}_n)$ (depends on the unknown distribution F) by $\text{Var}_{\hat{F}_n}(\hat{\theta}_n)$; if we still don't have a formula to compute this, we must do a further bootstrap estimate: we create B data sets x_{1i}, \dots, x_{ni} ($i = 1, \dots, B$) by drawing observations with replacement from the data, which simulates random variables $\sim \hat{F}_n$; we get estimates $\hat{\theta}_n(x_{1i}, \dots, x_{ni})$ for $i = 1, \dots, B$ of which we compute the sample variance.

Regular models, Fisher information Define the score function $U_\theta(x) := \frac{\partial}{\partial \theta} \log f(x, \theta) = f'_x(\theta)/f_x(\theta)$ (suppose likelihood function positive, differentiable and some other nice properties, so-called *regular models*). The Fisher Information of a model is the function $I: \theta \mapsto \text{Var}_\theta(U_\theta)$. This is used to approximate standard error of a MLE.

I vanishes on an interval iff f_x is constant on Θ_0 for almost all x , so no observation can distinguish between the parameters in Θ_0 . Additivity: an n -fold product model has Fisher information $I^{\otimes n} = nI$.

Asymptotic normality of MLE Under appropriate regularity conditions, $\sqrt{\text{Var}(\hat{\theta}_n)} \approx \sqrt{1/I_n(\theta)} =: \widehat{\text{SE}}$ and the statistic $(\hat{\theta}_n - \theta)/\widehat{\text{SE}}$ (or $(\hat{\theta}_n - \theta)/\text{SE}$) converges to $\mathcal{N}(0, 1)$ in distribution. Consequently we have the approximate $1 - \alpha$ confidence interval

$$C_n = \hat{\theta}_n \mp z_{\alpha/2} \widehat{\text{SE}}, \quad P_\theta(C_n \ni \theta) \rightarrow 1 - \alpha.$$

Delta method For $\tau = g(\theta)$ with g differentiable and $g'(\theta) \neq 0$, the MLE is $\hat{\tau} = g(\hat{\theta})$ (equivariance) and we have asymptotic normality as above, with $\text{SE}(\hat{\tau}_n) = |g'(\hat{\theta})| \widehat{\text{SE}}(\hat{\theta}_n)$.

Parametric bootstrap Alternatively, one can estimate standard errors using the bootstrap method sampling from $f(x, \hat{\theta}_n)$ (instead of \hat{F}_n as in the non-parametric bootstrap).

Cramér-Rao inequality (or information inequality). Repeating a regular experiment n times, the variance of an unbiased estimator of τ has order at least $1/n$, but there is a condition to get a best estimator, an estimator that is *Cramér-Rao efficient*:

Take a regular model, $\tau: \Theta \rightarrow \mathbb{R}$ a \mathcal{C}^1 function, $\tau' \neq 0$, T regular estimator of τ , i.e. satisfying

$$\frac{d}{d\theta} \text{E}_\theta(T) = \int T(x) \frac{d}{d\theta} f(x, \theta) dx.$$

Then

$$\text{Var}_\theta(T) \geq \tau'(\theta)^2 / I(\theta)$$

with equality iff $T - \tau(\theta) = \tau'(\theta)U_\theta/I(\theta)$ for all $\theta \in \Theta$, i.e.

$$f(x, \theta) = \exp[a(\theta)T(x) - b(\theta)]h(x)$$

for a primitive of I/τ' , $h: X \rightarrow (0, \infty)$ measurable function,

$$b(\theta) = \log \int_X \exp[a(\theta)T(x)]h(x)dx.$$

Exponential models Models with this likelihood function are called *exponential models* wrt the statistic T . Exponential models of one parameter have the increasing likelihood property which will be useful in tests. Examples: a family of

binomial distributions $\{B_{n,\theta} : \theta \in (0, 1)\}$ for fixed n , $T(x) = x/n$;
Poisson distributions $\{P_\theta : \theta > 0\}$, $T(x) = x$;
normal distributions $\{N_{\theta,v} : \theta \in \mathbb{R}\}$ for fixed variance v , $T(x) = x$;
the n -fold product of exponential models for $T(x)$, by setting $T_n(x) = \frac{1}{n} \sum_{i=1}^n T \circ X_i$.

Consistency A sequence of estimators $T_n(X_1, \dots, X_n)$ is *consistent* if $T_n \xrightarrow{P_\theta} \theta$ as $n \rightarrow \infty$ (this is analogous to the WLLN).

We restrict to the case of independent observations, so to the infinite product model $(X = E^{\mathbb{N}}, \mathcal{F} = \mathcal{E}^{\otimes \mathbb{N}}, P_\theta = Q_\theta^{\otimes \mathbb{N}} : \theta \in \Theta)$.

Consistency of \bar{X} and \hat{V} in real product models The sequences $(M_n)_{n \geq 1}$, $(V_n^*)_{n \geq 2}$ of sample mean and corrected variance are consistent.

Consistency of MLE's Consider a one-parameter standard model with Θ open interval, $Q_\theta \neq Q_{\theta'}$ for $\theta \neq \theta'$ (identifiable), and n -fold likelihood function $f^{\otimes n}(x, \theta) = \prod_{i=1}^n f(x_i, \theta)$ *unimodal* in θ for all x and $n \geq 1$, i.e. there is a MLE $T_n: E^{\mathbb{N}} \rightarrow \mathbb{R}$ such that the function is increasing for $\theta < T_n(x)$ and decreasing for $\theta > T_n(x)$. Then the sequence T_n is consistent.

Long proof, uses relative entropy.

Entropy The relative entropy (or Kullback-Leibler divergence) is

$$H(P, Q) := KL(P||Q) = E_P \left(\log \frac{f}{g} \right) = - \int p(x) \ln \left(\frac{q(x)}{p(x)} \right)$$

measures how distinguishable the two probability measures are (it's ≥ 0 with $=$ if $P = Q$). The entropy

$$H(X) = -E(\log p(X)) = - \int p(x) \ln p(x) dx$$

measures how "scattered" the distribution is (is it always well defined?), or the average amount of information given by a random variable X with density p . Given a joint density $p(x, y)$, the conditional entropy

$$H(y|x) = - \iint p(y, x) \ln p(y|x) dy dx$$

is the average additional information needed to specify y once x is known (because the additional information is $-\ln p(y|x)$ and satisfies $H(x, y) = H(y|x) + H(x)$). The mutual information $I(X, Y) = KL(p(X, Y) || p(X)p(Y))$ measures how far X and Y are from being independent, and satisfies $I(X, Y) = H(X) - H(X|Y) = H(Y) - H(Y|X)$.

2.1 Confidence regions

A *confidence region* with error level α (or confidence level $1 - \alpha$) for a characteristic $\tau: \Theta \rightarrow \Sigma$ to be estimated is a map $C: X \rightarrow \mathcal{P}(\Sigma)$ such that for all $\theta \in \Theta$

$$P_\theta(x \in X \mid \tau(\theta) \in C(x)) \geq 1 - \alpha.$$

(we assume $\{s \in C(-)\} \in \mathcal{F}$ for all $s \in \Sigma$). C is also called an *interval estimator* when $\Sigma = \mathbb{R}$ and every $C(x)$ is an interval.

Intuition ($\Sigma = \Theta$, $\tau = \text{id}$): given x , what θ 's are "most likely" to be the true one? Answer: those for which x is most likely to be observed (with probability at least $1 - \alpha$). These are contained in $C(x)$. In other words, the probability of observing an x (i.e. of generating an interval) such that $\theta \notin C(x)$ is less than α ; resampling in the long run, a percentage tending to $100(1 - \alpha)\%$ of the confidence intervals will contain the true value.

Constructing a confidence region for θ : for each $\theta \in \Theta$, define a set of most likely outcomes as $C_\theta := \{x \in X \mid f_\theta(x) \geq c_\theta\}$ for the smallest $c_\theta > 0$ needed to exceed the confidence level $1 - \alpha$: this determines how many x we should include in order to exceed the confidence level. Then set $C(x) := \{\theta \in \Theta \mid x \in C_\theta\}$.

In the prototypical case, the density f_θ is *unimodal*, i.e. it has a single peak: then C_θ is an interval around the peak, large enough to make sure the tails together have probability at most α .

Quantiles and fractiles Let $(\mathbb{R}, \mathcal{B}, P)$ be a probability space, $\alpha \in (0, 1)$. A number $q \in \mathbb{R}$ is called an α -*quantile* of P if $P((-\infty, q]) \geq \alpha$ and $P([q, \infty)) \geq 1 - \alpha$. The quantile of a real-valued random variable is defined as the quantile of its distribution.

When the data is divided in 4 groups we call them quartiles: a *lower quartile*, *median*, *upper quartile* is respectively a 1/4-, 1/2-, 3/4- quantile. When divided in 100 groups they are called percentiles: 25th percentile = 0.25-quantile = first quartile. An $(1 - \alpha)$ -quantile is also called an α -*fractile*.

Percentiles are used to summarize histograms that don't follow a normal distribution, and hence can't be satisfactorily summarized by giving average and SD. Conversely, if we are given a score (an observation), its percentile rank is the percentage of data below that score.

Look at the graphical representation of a quantile in the cumulative distribution function and the density function, respectively. One has:

$$\alpha = F_P(q) = P((-\infty, q]) = \int_{-\infty}^q f(x)dx$$

so q is a point where the value of the distribution crosses the level α , and the area subtended by the density in $(-\infty, q]$ is α . This q is unique if F_Q is strictly increasing in a neighborhood of q .

2.2 Statistical decision

We decompose $\Theta = \Theta_0 \sqcup \Theta_1$ where $\theta \in \Theta_0$ is the *null hypothesis* (the normal case) and $\theta \in \Theta_1$ is the *alternative* (the problematic case). If

Tests A test of Θ_0 against Θ_1 is a statistic $\varphi: X \rightarrow [0, 1]$ used to set a decision rule so that, on the basis of the observed x , we decide as follows: if $\varphi(x) = 0$ we accept the null hypothesis; if $\varphi(x) = 1$ we reject it; if $\varphi(x) \in (0, 1)$ we decide at random, with probability $\varphi(x)$ for rejection of the null hypothesis.

It is called *non-randomized* if $\varphi(x) \in \{0, 1\}$ for all x : then $R := \varphi^{-1}(1)$ is the *rejection region* or *critical region*, $\varphi^{-1}(0)$ is the *acceptance region*.

One-sided (left-sided) test: $H_0: \theta \leq \theta_0$. Two-sided test $H_0: \theta = \theta_0$ (simple hypothesis) i.e. a rejection region of the form $R = \{x: T(x) > c\}$ where T is a statistic whose distribution is known and a critical value c which is chosen so as to get the desired size α (below).

False rejection of H_0 is called type 1 error, or 'false positive'. False acceptance of H_0 is called type 2 error, or 'false negative'.

Significance level, size, power The *power function* of a test φ is

$$W_\varphi: \Theta \rightarrow [0, 1], \theta \mapsto E_\theta(\varphi) = P_\theta(X \in R)$$

- for $\theta \in \Theta_0$, $P_\theta(X \in R)$ is the probability of false rejection of H_0 . The *size* of φ is $\sup_{\theta \in \Theta_0} P_\theta(X \in R)$, the worst case probability of a type 1 error. α is used to denote either the size or, more generally, the level: if a test has size $\leq \alpha$ we say it is of significance level α .
- for $\theta \in \Theta_1$, $1 - \beta := P_\theta(X \in R)$ is the probability of true rejection of H_0 , called the *power* of φ at θ , where β denotes the probability of false acceptance of H_0 against θ . Usually we just test H_0 against a single alternative θ at a time, so we don't use a sup definition as before.

Example Random variable $X \sim N(\mu, \sigma^2)$, so that $\bar{X}_n \sim N(\mu, \sigma^2/n)$. Test for $H_0: \mu = \mu_0$ vs $H_1: \mu > \mu_0$. For example, if X is the difference between two treatments, we'll use $\mu_0 = 0$ to test whether the treatment is effective. The test statistic $Z_\mu = \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}}$ has standard normal distribution under H_0 , so we reject H_0 at the 0.05 significance level if $Z_{\mu_0} > q_0$, a $(1 - \alpha)$ -quantile of $N(0, 1)$, equivalently if $\bar{X} \geq \mu_0 + q_0 \frac{\sigma}{\sqrt{n}}$. The power at any μ_1 is

$$\begin{aligned} W(\mu_1) &= P(Z_{\mu_0} > q_0 | \mu = \mu_1) = P\left(\bar{X} \geq \mu + q_0 \frac{\sigma}{\sqrt{n}} \middle| \mu = \mu_1\right) = \\ &= P\left(Z_\mu \geq q_0 + \frac{\mu_0 - \mu}{\sigma/\sqrt{n}} \middle| \mu = \mu_1\right) = \\ &= P\left(Z_\mu \geq \frac{\mu_0 + q_0(\sigma/\sqrt{n}) - \mu}{\sigma/\sqrt{n}} \middle| \mu = \mu_1\right) \end{aligned}$$

The closer μ_0 and μ_1 are, the less the power.

A test φ should have small level and big power. A *uniformly most powerful (UMP) test* of level α is a test of level α whose power exceeds that of any other test of level α . The problem is finding the right balance between significance level and power of a test.

A test φ is *unbiased* of level α if, for all $\theta_0 \in \Theta_0$ and $\theta_1 \in \Theta_1$, $W_\varphi(\theta_0) \leq \alpha \leq W_\varphi(\theta_1)$ (one decides for the alternative with a higher probability when it is right than when it is false).

Relation between confidence regions and tests For $\theta_0 \in \Theta$, an α -confidence region $\{\theta_0 \notin C(-)\}$ is the rejection region of a test for $H_0: \theta = \theta_0$ against $H_1: \theta \neq \theta_0$, and conversely, given a such test for every $\theta_0 \in \Theta$, they can be combined to construct an α -confidence region.

Terminology in binary classification

$$\begin{aligned} P + N &= (TP + FN) + (TN + FP) = \\ &= (TP + FP) + (TN + FN) = \text{class}(P) + \text{class}(N) \end{aligned}$$

The *confusion matrix* for a binary classifier (easily generalized to multiple classes) is

		Prediction		Total
		+	−	
True status	+	TP	FN	P
	−	FP	TN	N
Total		$\text{class}(P)$	$\text{class}(N)$	total

- type 1 error = false positive;
type 2 error = false negative
- false positive rate = fall-out = FP/N (corresponds to the significance level, which is however an a-priori property of the test)
specificity = true negative rate TN/N (its complement).
- sensitivity = recall = true positive rate = TP/P (corresponds to the statistical power of the test),
false negative rate = FN/P .
- positive predictive value = precision = $TP/\text{class}(P)$
false discovery rate = $FP/\text{class}(P)$
- negative predictive value = $TN/\text{class}(N)$
false omission rate = $FN/\text{class}(N)$
- prevalence = $P/\text{total} = P/(P + N)$;
accuracy = $(TP + TN)/(P + N)$

If prevalence is low we may have high specificity but low positive predictive value. With skewed datasets (one class has much higher frequency than the others) high accuracy is not a good metric.

Precision-recall tradeoff To compare precision and recall between different models (different thresholds) you can plot the PR curve with graph (recall, precision) or the ROC curve with graph (FPR, TPR=recall). The AUC score (area under curve) can then be used to choose the best classifier.

The F_1 score, harmonic mean of precision and recall, is only high when both values are high:

$$\frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Experiment procedure First we choose a level α and the type of test φ ; φ may depend on a parameter $c \in X$ (e.g. the number of times something should be guessed correctly for us to conclude extrasensory power), which is chosen so that $\beta_\varphi(\theta) \leq \alpha$ for all $\theta \in \Theta_0$.

The threshold c (hence, the size) has to be chosen before the experiment. If we modify it after the experiment we are essentially using a test $\psi = 1$ identically (unless we repeat the experiment...).

Note that depending on the probability P_θ involved, certain thresholds may give a better power function even if they require a lower percentage of successes.

To avoid fixing a threshold, we use p-values:

p-value

- Abstract definition: a statistic with $p(X) \in [0, 1]$ and $\sup_{\theta \in \Theta_0} P_\theta(p(X) \leq \alpha) \leq \alpha$. Given this, we can construct a level α test as a test that rejects H_0 iff $p(X) \leq \alpha$.
- If we are given tests T_α of level $\alpha = \sup_{\theta \in \Theta_0} P_\theta(X \in R_\alpha)$ for every $\alpha \in (0, 1)$ with rejection regions R_α , then $p(x) = \inf\{\alpha \mid T_\alpha(x) \in R_\alpha\}$ is the smallest level at which H_0 is rejected.
- Most common case: $\{T > c\}$, $p(x) = \sup_{\theta \in \Theta_0} P_\theta(T \geq T(x))$. Probability of observing $T \geq T(x)$ under the null hypothesis: the smaller it is, the stronger the evidence for H_1 .

Don't confuse it with the test size: p-value is a property of the result, whereas size is a property of the test. A result is statistically significant when the p-value is $\leq \alpha$.

p -hacking or data snooping is the practice of testing multiple hypotheses on a single data set in order to find a statistically significant result which is in fact a spurious correlation (false positive, or false rejection).

Adjusted p-values To avoid the problem of spurious significance when testing multiple hypotheses, we must decrease the p-values according to some rule. We must control the family-wise error rate, i.e. the probability of making at least one type 1 error, given by

$$FWER = P(\text{\#type 1 errors} \geq 1) = 1 - (1 - \alpha)^m$$

where m is the number of independent tests. If tests are not independent, we still have $FWER \leq m\alpha$ by subadditivity of probability measures. Usual corrections to solve this: Bonferroni correction replaces α with α/m , to keep

the FWER below α , but raising the risk of false acceptance of null hypothesis; Holm-Bonferroni sorts the p-values p_1, \dots, p_m in increasing order and replaces α_i with $\alpha/(m - i + 1)$, which controls FWER as before but causes a smaller increase in type 2 error risk compared to the previous method.

See q-value.

When multiple comparisons are performed, we define the false positive rate as $\frac{\# \text{type 1 errors}}{\text{true null hypotheses}}$.

Effect size Can be given for example as:

- Pearson's r
- Cohen's d: $\frac{\bar{X}_1 - \bar{X}_2}{\text{pooled } s \text{ or } \sigma}$ measuring how much the two distributions overlap, by combining mean and sd in a single metric.
- odds ratio.

Power analysis Given 3 of the following, the remaining quantity is determined: effect size, sample size, test size, power. Power analysis ensures that we choose the correct sample size to make sure that we have enough data to make a good decision when we reject or accept H_0 , and we avoid resorting to "p-hacking". Remember the larger the sample size, the lower the standard deviation around the mean, the less overlap between two distributions... we use previous data or educated guesses to choose the means and sd data needed for power analysis calculations.

Effect size (usually standardized) is for example the difference between two means; it is not a strictly statistical measure, but a more practical one. Increasing sample size increases power: power should not be too high otherwise the test will be overly sensitive and any correlation will become significant. Usual power level is about 80%.

Neyman-Pearson test We can find the most powerful test of level α for a standard model of the form $(X, \mathcal{F}; P_0, P_1)$. A Neyman-Pearson test has the form $\varphi(x) = 1$ if $R(x) := f_1(x)/f_0(x) > c$ (*likelihood ratio*) for some threshold c , and 0 if $R(x) < c$. Choose c an α -fractile of $P_0 \circ R^{-1}$. What value if $R(x) = c$? We want α to equal $E_0(\varphi) = P_0(R > c) + \gamma P_0(R = c)$ for some γ .

$$\begin{aligned} P_0(R = c) &= -P_0(R > c) + P_0(R \geq c) \\ &\geq -P_0(R > c) + \alpha \end{aligned}$$

If $P_0(R = c) = 0$ then $P_0(R > c) = \alpha$ and we can take $\varphi(x) = 0$, otherwise take $\gamma := (\alpha - P_0(R > c))/P_0(R = c)$. Every such Neyman-Pearson test is a most powerful test of level α , and every most powerful test of this level is indistinguishable from a Neyman-Pearson test.

How quickly does the power increase for independently repeated observations? Consider the infinite product model $(X = E^{\mathbb{N}}, \mathcal{F} = \mathcal{E}^{\otimes \mathbb{N}}, P_\theta = Q_\theta^{\otimes \mathbb{N}} : \theta \in \{0, 1\})$ associated to a simple model as above, and φ_n be size- α Neyman-Pearson tests using only observations X_1, \dots, X_n , where X_i are the projections. Then $E_1(\varphi_n) \approx 1 - e^{-n} H(Q_0, Q_1)$, more precisely

$$\lim_n \frac{1}{n} \log(1 - E_1(\varphi_n)) = -H(Q_0, Q_1)$$

One-sided tests Neyman-Pearson can be generalized for models with increasing likelihood ratios wrt a statistic $T: X \rightarrow \mathbb{R}$, i.e. $T(x) < T(y)$ implies $R_{\theta':\theta}(x) < R_{\theta':\theta}(y)$. We can test the one-sided (left-sided) test problem $H_0: \theta \leq \theta_0$ with a UMPT φ of level α of the form

$$\varphi(x) = \begin{cases} 1 & T(x) > c \\ \gamma & T(x) = c \\ 0 & T(x) < c \end{cases}$$

with c and γ uniquely determined by the condition $G_\varphi(\theta_0) = \alpha$. Moreover, G_φ is increasing.

Construct φ as in the simple case. To show it's a UMPT we use the increasing property just mentioned: take some $\theta < \theta'$, then $R_{\theta':\theta}(x) < R_{\theta':\theta}(c)$ implies $T(x) < T(c)$ hence $\varphi(x) = 0$ and $>$ implies $\varphi(x) = 1$, so φ is a N-P test of the simple null hypothesis $\{\theta\}$ against $\{\theta'\}$; for $\theta = \theta_0 < \theta' \in \Theta_1$ this shows it's a UMPT. To show the level is α , we should look at all other $\theta \in \Theta_0$ too; it's enough to show G_φ is increasing. Since the simple test of θ against θ' is more powerful than the constant test ψ , then $G_\varphi(\theta') > G_\psi(\theta') = G_\varphi(\theta)$ and we conclude.

The right-sided hypothesis can be tested multiplying both T and θ by -1 .

2.3 Tests for Gaussian models

General procedure: we measure by how many SD's the observed value of the statistic (score) differs from a benchmark (e.g. μ_0 , chosen depending on the null hypothesis) and calculate the p-value; if it's low enough, we can reject the null hypothesis.

Paired tests: used for paired data, consisting of groups which are not independent. E.g. a measurement on the *same* patient before and after treatment.

First some useful distributions that will show up.

χ , F , t -distributions Let $X_i \sim N_{0,1} \sim Y_j$ iid. Then:

- The chi-square χ_n^2 is $\Gamma_{1/2, n/2}$, which is the distribution of $\sum_{i=1}^n X_i^2$.
- Fisher's $F_{m,n}$ -distribution is the distribution of $(\frac{1}{m} \sum_1^m X_i^2) / (\frac{1}{n} \sum_1^n Y_i^2)$, i.e. the ratio $(S_m/m) / (S_n/n)$ where the S_i are independent χ_i^2 distributions.
- Student's t_n -distribution is the distribution t_n of $T = X_1 / \sqrt{\frac{1}{n} \sum_{j=1}^n Y_j^2}$. Note that $T^2 \sim F_{1,n}$.

n -fold Gaussian product model $(\mathbb{R}^n, \mathcal{B}^n, N_{\mu,v}^{\otimes n}: \mu \in \mathbb{R}, v > 0)$. Student's theorem: M and V^* are independent and:

- $M \sim N_{\mu, v/n}$, or equivalently $(M - \mu)\sqrt{n}/\sigma \sim N_{0,1}$;
- $\frac{n-1}{v} V^* = \frac{1}{\sigma^2} \sum_1^n (X_i - \bar{X})^2 \sim \chi_{n-1}^2$

- $(M - \mu)\sqrt{n/V^*} \sim t_{n-1}$ and an $(1 - \alpha)$ -confidence interval for μ is

$$(M \mp t\sqrt{V^*/n}), \quad t = \alpha/2\text{-fractile of } t_{n-1}.$$

The 95% Gaussian CI is given by $t\sqrt{V^*} = t\sigma^*$ where t is the $1 - \alpha/2 = 0.975$ -quantile of $N_{0,1}$ i.e. $t = \Phi^{-1}(0.975) = 1.96$.

The following are UMP tests of level α .

Z-test The Z -score or standard score of an observation (raw score) x is $(x - E(X))/\sigma(X)$, the distance measured in standard deviations of x from the population mean (this means converting the value to standard units). The p -values are $\Phi(Z)$, $\Phi(-Z)$, $2\Phi(-|Z|)$ for left-, right-, two-sided tests respectively.

- One-sided χ^2 -test of $H_0: v \leq v_0$ with rejection region

$$\sum_{i=1}^n (X_i - \bar{X})^2 > v_0 \chi_{n-1, 1-\alpha}^2$$

is a UMP test of level α . If μ is known replace \bar{X} by μ ; then the statistic is $\sim \chi_n^2$ (so we should use $\chi_{n, 1-\alpha}^2$) and the product model is exponential wrt it, so by the above theory we get a UMPT.

If μ is unknown we replace it with \bar{X} and by the above theory the $nV/v \sim \chi_{n-1}^2$ so the one in the statement is a test of level α . To check that it's a most powerful test is more complicated.

The symmetric statement does not hold: a UMPT does not exist in this case; the test of $H_0: v \geq v_0$ with rejection region

$$\sum_{i=1}^n (X_i - \bar{X})^2 < v_0 \chi_{n-1, \alpha}^2$$

is a UMP test only within the class of *unbiased* level- α tests.

- One-sided Student's t -test of $H_0: \mu \leq \mu_0$ with rejection region

$$T(X) := \frac{\bar{X} - \mu_0}{S^*/\sqrt{n}} > t_{n-1, 1-\alpha}$$

(the α -fractile of the t_{n-1} distribution) is a UMP test within the class of unbiased level- α tests (a UMPT does not exist). Invert inequalities for the opposite test.

One-sided Gauss test: if σ is known, we can replace S^* by σ and the statistic has a $\mathcal{N}(0, 1)$ distribution so we replace $t_{n-1, 1-\alpha}$ by $\Phi^{-1}(1 - \alpha)$. This is a UMP test.

- Two-sided Student's t -test of $H_0: \mu = \mu_0$ with rejection region

$$\frac{|\bar{X} - \mu_0|}{S^*/\sqrt{n}} > t_{n-1, 1-\alpha/2}$$

is a UMP test within the class of unbiased level- α tests. If σ is known this is just the Wald test below.

Two-sample tests The two-sample Z-test is used to test the difference between two means, $H_0: \mu_1 - \mu_2 = d$ (usually 0) when the variances are known.

$$\frac{\bar{X}_1 - \bar{X}_2 - (\mu_1 - \mu_2)}{\sigma(\bar{X}_1 - \bar{X}_2)} \sim N(0, 1)$$

where $\sigma(\bar{X}_1 - \bar{X}_2)$ is, depending on whether the test is independent or paired, respectively

$$\sqrt{\frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2}} \quad \text{or} \quad \sqrt{\frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2} - 2\text{Cov}(\bar{X}_1, \bar{X}_2)}$$

In the paired case, if the data consists of pairs (X_{i1}, X_{i2}) for $i = 1 \dots, n$ and $D_i := X_{i1} - X_{i2}$, then we are interested in

$$\frac{\bar{D} - d}{\sigma(\bar{D})} \tag{2.1}$$

which is the same statistic as before since

$$\bar{D} = \frac{1}{n} \sum_{i=1}^n (X_{i1} - X_{i2}) = \bar{X}_1 - \bar{X}_2;$$

in this case $\text{Cov}(\bar{X}_1, \bar{X}_2) = \text{Cov}(X_{i1}, X_{i2})/n$ is the within-pair (or within-subject) correlation.

When variances are unknown we use t-tests; the paired ones are easily done using paired data, replacing σ with the sample standard deviation:

$$\frac{\bar{D} - d}{s(\bar{D})} \sim t_{n-1}, \quad s(\bar{D}) = \frac{1}{n-1} \sum_{i=1}^n (D_i - \bar{D})^2$$

so we can focus on the independent tests. When the two variances can be assumed to be the same, we use

$$\frac{\bar{X}_1 - \bar{X}_2 - (\mu_1 - \mu_2)}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \sim t_{n_1+n_2-2}$$

where s_p is the pooled sample standard deviation, the unbiased estimate of the common standard deviation, which has $n_1 + n_2 - 2$ DoF. When the two variances can't be assumed to be equal, we use the Welch t-test:

$$\frac{\bar{X}_1 - \bar{X}_2 - (\mu_1 - \mu_2)}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$

whose distribution is approximately a t with number of DoF calculated by the Welch-Satterthwaite equation:

$$\frac{\left(\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}\right)^2}{\frac{(s_1^2/n_1)^2}{n_1-1} + \frac{(s_2^2/n_2)^2}{n_2-1}}$$

Multiple testing We have m tests of hypotheses H_{0i} with p-values P_i , $i = 1, \dots, m$. The chance of at least one false rejection is higher than that of each single test. The Bonferroni method is to reject H_{0i} if $P_i < \alpha/m$; then the probability of falsely rejecting any null hypothesis is at most α (by subadditivity of probability).

2.4 Goodness of fit and independence tests

The following 3 tests are equivalent under H_0 and are used to compare a full model against a simpler nested model which imposes some constraints on the parameters.

Likelihood ratio test for $H_0: \theta \in \Theta_0$; the likelihood ratio statistic is

$$\lambda = 2 \log \left(\frac{\sup_{\Theta} \mathcal{L}(\theta)}{\sup_{\Theta_0} \mathcal{L}(\theta)} \right) = 2 \log \left(\frac{\mathcal{L}(\hat{\theta})}{\mathcal{L}(\hat{\theta}_0)} \right)$$

(hats denote MLE's). Wilks' theorem states that $\lambda \rightsquigarrow \chi_{r-q, \alpha}^2$ where $r = \dim \Theta$ and $q = \dim \Theta_0$.

Wald test The Wald statistic $(\hat{\theta} - \theta_0)^2 / \text{Var}(\hat{\theta})$ converges asymptotically to χ_1^2 under H_0 . For asymptotically normal $\hat{\theta}$ the Wald test of size α (asymptotically) of $H_0: \theta = \theta_0$ rejects when $|(\hat{\theta} - \theta_0) / \widehat{\text{SE}}| > z_{\alpha/2}$. Equivalently we are testing whether θ_0 is in the asymptotic $1 - \alpha$ confidence interval $\hat{\theta} \mp \widehat{\text{SE}} z_{\alpha/2}$. If the true value is $\theta_* \neq \theta_0$ the power $\beta(\theta_*)$ is approximately... We approximate the p-value using Φ as above.

Score test $U(\theta_0)^2 / I(\theta_0)$ has an asymptotic distribution χ_1 under H_0 . Generalizes to k parameters using χ_k .

Empirical distribution The absolute frequency $h_n(i) = \#\{k \in \{1, \dots, n\} \mid X_k = i\}$; the vector of relative frequencies $L_n = (\frac{h_n(1)}{n}, \dots, \frac{h_n(s)}{n})$ is the *histogram*. The *empirical distribution* is $L = \frac{1}{n} \sum_{i=1}^n \delta_{X_i}$ and it's the unique MLE of θ ; empirical distribution function $\hat{F}_n: x \mapsto \frac{1}{n} \sum_{i=1}^n I(X_i \leq x)$. The quantiles of L are called *sample quantiles* (order statistics below are a special case). The Glivenko-Cantelli theorem states that $\sup_x |\hat{F}_n(x) - F(x)|$ converges to 0 almost surely. (note: histograms as a representation can also have bars with variable width where vertical axis is frequency density and the bars have area proportional to frequency).

A statistical functional is any statistic $T(F)$. Its plug-in estimator is $T(\hat{F}_n)$. For a linear functional $\int r(x) dF(x)$ the plug-in estimator becomes $\frac{1}{n} \sum r(X_i)$.

Likelihood-ratio goodness-of-fit test Test of the null hypothesis $H_0: \theta = f$ to test, after n trials, if the observed distribution is the one we expect. The test should reject for big enough values of the log-likelihood ratio

$$\log \frac{\sup_{\theta \neq f} \prod_{i=1}^s \theta(i)^{h_n(i)}}{\prod_{i=1}^s f(i)^{h_n(i)}} = \log \frac{\sup_{\theta \in \Theta} \dots}{\dots} = n \sum_{i=1}^s L_n(i) \log \frac{L_n(i)}{f(i)} = nH(L_n, f)$$

which has the quadratic approximation

$$2nH(L_n, f) \xrightarrow{P_f} D_{n,f} := n \sum_{i=1}^s \frac{L_n(i)^2}{f(i)} - n$$

and under H_0 this $\rightsquigarrow \chi_{s-1}^2$ i.e. $\lim_n P_f(D_{n,f} \leq c) = \chi_{s-1}^2([0, c])$. So the test with rejection region

$$D_{n,f} > \chi_{s-1, 1-\alpha}^2$$

has approximately size α for n large enough (for small n one should use the exact distribution of $D_{n,f}$ which can be derived from the multinomial distribution). (write it in terms of p-value?)

Pearson's χ^2 multinomial test $X = (X_1, \dots, X_k) \sim \text{multinomial}(n, \mathbf{p})$, the MLE of \mathbf{p} is $\hat{\mathbf{p}} = (X_1/n, \dots, X_k/n)$. We want to test $H_0: \mathbf{p} = (p_1, \dots, p_k)$. We define the statistic

$$T(X) = \sum_{j=1}^k \frac{(X_j - np_j)^2}{np_j}$$

and, under H_0 we have $np_j = E(X_j)$ and $T \rightsquigarrow \chi_{k-1}^2$ so the test rejecting if $T > \chi_{k-1, \alpha}^2$ has asymptotic level α .

Classic example: test if a die is fair, using $\mathbf{p} = (n/6, \dots, n/6)$.

Contingency tables A contingency table (or cross tabulation) records $h_n(ij)$ for $ij := (i, j) \in E = A \times B = \{1, \dots, a\} \times \{1, \dots, b\}$.

	$B = 1$	$B = 2$	\dots	$B = b$	
$A = 1$	h_{11}	h_{12}	\dots	h_{1b}	$h_{1\bullet}$
$A = 2$	h_{21}	h_{22}	\dots	h_{2b}	$h_{2\bullet}$
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
$A = a$	h_{a1}	h_{a2}	\dots	h_{ab}	$h_{a\bullet}$
	$h_{\bullet 1}$	$h_{\bullet 2}$	\dots	$h_{\bullet b}$	n

Alternatively it can be written in terms of probabilities $p_{ij} := h_{ij}/n$. A typical contingency table is one where variable A takes values {diseased, healthy} and B is {exposed, not exposed} (e.g. to a pathogen).

Consider all strictly positive densities on E

$$\Theta = \left\{ \theta = \theta(ij)_{ij \in E} \in (0, 1)^E : \sum_{ij \in E} \theta(ij) = 1 \right\}$$

The null hypothesis of independence of variables A and B can be written as $H_0: \theta = \theta^A \otimes \theta^B$ product of the marginal distributions, $\theta^A(i) = \sum_{j \in B} \theta(ij)$ (denote Θ_A the set of marginal distributions on A ...), so

$$\Theta_0 = \left\{ \alpha \otimes \beta = (\alpha(i)\beta(j))_{ij \in E} : \alpha \in \Theta_A, \beta \in \Theta_B \right\}$$

Pearson's χ^2 test of independence The null hypothesis H_0 is the independence of two variables in a contingency table; under H_0 we have

$$\sum_{i=1}^I \sum_{j=1}^J \frac{(X_{ij} - E_{ij})^2}{E_{ij}} \rightsquigarrow \chi_{(I-1)(J-1)}^2$$

with $E_{ij} = n\hat{p}_{ij} = n\hat{p}_{i\bullet}\hat{p}_{\bullet j} = \frac{X_{i\bullet}X_{\bullet j}}{n}$, the expected number of observations in cell ij .

Likelihood-ratio test of independence The empirical joint distribution is $L_n = (L_n(ij))_{ij \in E} = (h_n(ij)/n)_{ij \in E}$. Accept H_0 if L_n is sufficiently close to $L_n^A \otimes L_n^B$, with the distance measured similarly as before computing the log-likelihood ratio

$$\sum_{ij \in E} h_n(ij) \log \left(\frac{L(ij)}{L^A(i)L^B(j)} \right) = nH(L, L^A \otimes L^B)$$

and finding an analogous statistic \tilde{D}_n (as for the goodness-of-fit test) which $\rightsquigarrow \chi_{(a-1)(b-1)}^2$ under P_f , so a test with approximate size α is the one with rejection region

$$\tilde{D}_n > \chi_{(a-1)(b-1); 1-\alpha}^2$$

.

G test likelihood ratio test replacing χ^2 ...

Kolmogorov-Smirnov test of goodness of fit One-sample K-S test for H_0 : (the sample comes from F); statistic $D_n = \sup_x |\hat{F}_n(x) - F(x)|$; under H_0 then $D_n \rightarrow 0$ almost surely (Glivenko-Cantelli theorem) and $\sqrt{n}D_n \rightsquigarrow \sup_t |B(F(t))|$; the Kolmogorov distribution is that of

$$K = \sup_{t \in [0,1]} |B(t)| \quad \text{with cdf} \quad F(x) = 1 - 2 \sum_{k=1}^{\infty} (-1)^{k-1} \exp(-2k^2 x^2)$$

so the test rejects at level α if $\sqrt{n}D_n > K_{1-\alpha}$ where $P(K \leq K_{1-\alpha}) = 1 - \alpha$.

Two-sample K-S test for $H_0: F_1 = F_2$. $D_{n,m} = \sup_x |\hat{F}_{1,n}(x) - \hat{F}_{2,m}(x)|$, we reject when $\sqrt{\frac{nm}{n+m}} D_{n,m} > K_{1-\alpha}$.

2.5 Non-parametric methods

We don't assume a particular class of distributions.

Fisher's exact test With small samples, instead of the above χ^2 test of independence, we use this exact test. Above situation with $A = B = \{1, 2\}$ (e.g. positive or negative outcome of medical test).

	$B = 1$	$B = 2$	
$A = 1$	h_{11}	h_{12}	a_1
$A = 2$	h_{21}	h_{22}	a_2
	b_1	b_2	n

Under the null hypothesis $\theta \in \Theta_0$ of independence (probability of being in A1 or A2 does not depend on whether we are in B1 or B2) we have

$$\begin{aligned} P_\theta(h_{n11} = h_{11} \mid h_n^A 1 = a_1, h_n^B 1 = b_1) &= \\ &= \mathcal{H}_{b_1; a_1, a_2}(\{h_{11}\}) = \frac{\binom{a_1}{h_{11}} \binom{a_2}{h_{12}}}{\binom{n}{b_1}} = \frac{\binom{b_1}{h_{11}} \binom{b_2}{h_{21}}}{\binom{n}{a_1}} = \mathcal{H}_{a_1; b_1, b_2}(\{h_{11}\}) \end{aligned}$$

(the probability that h_{11} elements are positive in a random sample, without replacement, of b_1 elements chosen from a set of n of which a_1 are positive: it's a hypergeometric distribution).

Permutation test $X_1, \dots, X_m \sim F_X$ and $Y_1, \dots, Y_n \sim F_Y$, we want to test $H_0: F_X = F_Y$. Consider $T(X_1, \dots, X_m, Y_1, \dots, Y_n) = |\bar{X}_m - \bar{Y}_n|$. Compute T_1, \dots, T_B by applying $B \leq N!$ different permutations to the arguments of T : under the null hypothesis they are uniformly distributed, we compute the approximate p-value $P_0(T > t) = \frac{1}{B} \sum_{i=1}^B I(T_i > t)$.

Order and rank tests We only make use of the natural ordering of \mathbb{R} . Take iid $X_1, \dots, X_n \sim Q$ unknown but assumed continuous, i.e. $Q(\{x\}) = 0$ for all $x \in \mathbb{R}$, equiv. F_Q continuous (true if Q has a density). This ensures that there are no ties, i.e. that $P(X_i = X_j) = 0$ for $X_i \neq X_j$.

The *order statistics* $X_{1:n}, \dots, X_{n:n}$ are the X_i put in increasing order. If there are ties they can be defined more generally as $X_{j:n} = \min\{c \in \mathbb{R} \mid \sum_{i=1}^n 1_{X_i \leq c} \geq j\}$ (the smallest number exceeding at least j of the other outcomes, the smallest j/n quantile of the empirical distribution L ; the collection of order statistics is in 1-1 correspondence with L).

The median only depends on the ordering: if μ a median of Q , $T: \mathbb{R} \rightarrow \mathbb{R}$ order-preserving, then $T(\mu)$ is a median of $Q \circ T^{-1}$.

Rank statistics of a sample sequence X_1, \dots, X_n are R_1, \dots, R_n given by $R_i = \#\{j \mid X_j \leq X_i\}$ is the position of X_i in the ordered statistics, hence $X_i = X_{R_i:n}$.

CI for the median Let $b_n(\alpha)$ be the largest α -quantile of the binomial distribution $B_{n, 1/2}$; then

$$[X_{k:n}, X_{n-k+1:n}] \quad \text{with } k = b_n(\alpha/2)$$

is a confidence interval for $\mu(Q)$ of level α . The *sample median* is

$$\mu(L) = \begin{cases} X_{k+1:n} & n = 2k + 1 \\ \frac{X_{k:n} + X_{k+1:n}}{2} & n = 2k \end{cases}$$

Sign test for the median iid $X_1, \dots, X_n \sim Q$ continuous, $\mu_0 \in \mathbb{R}$. A test of level α for the two-sided problem $H_0: \mu(Q) = \mu_0$, the acceptance region $\{X_{k:n} \leq \mu_0 \leq X_{n-k+1:n}\}$ with $k := b_n(\alpha/2)$;

for the one-sided problem $H_0: \mu(Q) \leq \mu_0$, the acceptance region $\{X_{k:n} \leq \mu_0\}$ with $k := b_n(\alpha/2)$ whose power function is a strictly increasing function of $p(Q) := Q(\mu_0, \infty)$.

One-parameter exponential family with strictly increasing statistic and coefficient function: then $P_\theta \prec P_{\theta'}$ for $\theta < \theta'$.

Exact test for $H_0: \theta = \theta_0$ has the form

$$\text{p-value} = p = \sum_{y: T(y) \geq T(x)} P_{\theta_0}(X = y)$$

the probability of observing any outcome more extreme than x for our statistic T , under H_0 . It is in contrast with an approximate test like the χ^2 , where the statistic only asymptotically approaches a certain distribution.

Binomial test one-tailed: $H_0: \theta = \theta_0$, $H_1: \theta < \theta_0$.

$$p = \sum_{i=0}^k P(X = i) = \sum_{i=0}^k \binom{n}{i} \theta_0^i (1 - \theta_0)^{n-i}$$

One-tailed with $H_1: \theta > \theta_0$ is analogous with summation $\sum_{i=k}^n$. For the two-tailed the summation set is slightly more complicated.

In general, multinomial test for a multinomial distribution. For larger sample sizes, it is approximated with a likelihood ratio test or Pearson's chi-squared test.

Sign test special case of binomial test, with $\theta_0 = 1/2$. Two-sample test to determine if two medians are significantly different (tests the sign of the median of $X - Y$); one-sample test to determine if a median is significantly different from a given value.

Alternatively: paired t-test, Wilcoxon signed-rank test are more powerful because they make use of the magnitude rather than just the sign.

Median test for unpaired samples: compute the overall median with the joined groups, make a contingency table with A1 and A2 being number of samples above and below median respectively, and B1 and B2 the two groups to compare. Then use a Fisher exact test or a chi-squared with this contingency table.

Wilcoxon signed-rank test (or Wilcoxon T-test) the one-sample version tests if the median of X is nonzero: assign ranks R_i to the samples $|X_j|$ and define the statistic $T = \sum \text{sgn}(X_i) R_i$. For small n , the distribution of T under H_0 is found computing the 2^n possible sign combinations, and the probability of $T = t$ is obviously the number of combinations giving $T = t$ divided by 2^n .

The paired version follows by replacing X_i with $X_i - Y_i$.

Mann-Whitney U test (or Wilcoxon rank-sum test) is the analogue of the sign and Wilcoxon tests but for independent samples. Suppose we have $n = k + l$ observations with $X_1, \dots, X_k \sim P$ and $X_{k+1}, \dots, X_n \sim Q$. We want to test $H_0: P = Q$ against $H_1: P \prec Q$ where $P \preceq Q$ means stochastic dominance, i.e. $P(c, \infty) \leq Q(c, \infty)$ for all $c \in \mathbb{R}$. This can be used as a test for a difference in medians, i.e. to check if one group has a significantly higher median than the other.

Set $S_P := R_1 + \dots + R_k$ and $S_Q := R_{k+1} + \dots + R_{k+l}$. If $S_P < S_Q$ the majority of outcomes of the P -sample is further to the left of the real axis. How much smaller should it be to conclude $P \prec Q$? Sufficient to consider $S := S_P$ since $S_P + S_Q = n(n+1)/2$. We're interested in the U statistic:

$$S = U + \frac{k(k+1)}{2}, \quad U = U_{k,l} := \sum_{i=1}^k \sum_{j=k+1}^{k+l} 1_{X_i > X_j}$$

which can also be defined as

$$U = \min_{i=1,2} \left(n_1 n_2 + \frac{n_i(n_i+1)}{2} - S_i \right).$$

Then the test has rejection region $\{U < c\} = \{W < c + k(k+1)/2\}$, where U , under the hypotheses $P = Q$ continuous, has the distribution:

$$P^{\otimes n}(U = m) = N(m; k, l) / \binom{n}{k}$$

where N is the number of partitions of $\sum_{i=1}^k m_i = m$ into $m_1 \leq m_2 \leq \dots \leq m_k$ from the set $\{0, \dots, l\}$. In particular $P^{\otimes n}(U = m) = P^{\otimes n}(U = n - m)$.

$N(m; k, l)$ can be computed recursively as $\sum_{j=0}^k N(m - j; j, l - 1)$.

For large values of k, l , a normal approximation exists.

The corresponding two-sided problem has rejection region $\{U_P < c\} \cup \{U_Q < c\}$.

Kruskal-Wallis extends Mann-Whitney beyond 2 groups, in the same way that ANOVA is an extension of the t-test in the parametric world.

Histogram method Out of N observations, n_i of them fall into bin i of width Δ_i ; then we estimate $p_i = n_i / N \Delta_i$. The width determines the smoothness of the model (from biased to overfitting). Computational advantage: the data set can be discarded once the histogram has been computed. Issues: the estimated density is discontinuous; curse of dimensionality: for M bins and D variables we have M^D total bins.

General idea: an idea of local neighborhood based on some distance (here Euclidean) and a smoothing parameter. We start from an approximation $p(\mathbf{x}) = K/NV$ of the probability density of \mathbf{x} falling into a region of volume V containing K points. Fixing V will give kernel methods, fixing K will give nearest-neighbor methods.

Note: in the following methods the training set must be stored in order to be also used during the prediction phase (memory-based methods).

Kernel methods we estimate the density at \mathbf{x} with

$$p(\mathbf{x}) = \frac{1}{N} \sum_{n=1}^N K_\lambda(\mathbf{x} - \mathbf{x}_n)$$

where K_λ is a kernel function with bandwidth λ ; for example, the basic idea is to use $K_\lambda = \mathbb{1}_{[-\lambda/2, \lambda/2]}/\lambda^D$ so that we get the number of data points inside a cube of side λ centered at \mathbf{x} , divided by the region volume. However, this is still discontinuous, so we can choose smoother kernels, as long as they satisfy $K(\mathbf{u}) \geq 0$ and $\int K(\mathbf{u})d\mathbf{u} = 1$. The Gaussian kernel is

$$K_\sigma(\mathbf{u}) = \varphi_{\mathbf{u}, \sigma^2}$$

Issue: we want to be able to change λ depending on the data density of the region.

Nearest-neighbors Now we fix the parameter K and let V be the volume of the sphere containing exactly K points.

It can be extended to classification: the estimates $p(\mathbf{x}|C_k) = K_k/N_k V$, $p(\mathbf{x}) = K/NV$, $p(C_k) = N_k/N$, and using Bayes' theorem the posterior probability is

$$p(C_k|\mathbf{x}) = \frac{p(\mathbf{x}|C_k)p(C_k)}{p(\mathbf{x})} = \frac{K_k}{K}$$

so the classifier assigns to the class with highest probability.

2.6 Epidemiology

	Event	Non-Event	
Control	CE	CN	C
Intervention	IE	IN	I
	E	N	n

Control/Intervention can be Exposed/Non-Exposed, and Event/Non-Event can be Diseased/Healthy.

- Event rates for each group:
intervention event rate IE/I , estimating $P(\text{disease}|\text{exposed})$.
control event rate CE/C which estimates $P(\text{disease}|\text{unexposed})$, the baseline risk. The term incidence (in a given group) may refer to either the number or the rate of events in that group. Prevalence is $P(\text{disease})$, the rate of the disease in the population, whereas incidence has to do with the occurrence of new cases.
- Relative risk (or risk ratio) $RR = \frac{IE/I}{CE/C}$
- Absolute risk difference (reduction or increase, depending on the sign), is $CE/C - IE/I$. Its inverse is the number needed to treat (or harm).
- relative risk difference $\frac{CE/C - IE/I}{CE/C} = 1 - RR$.
- Odds ratio $OR = \frac{IE/IN}{CE/CN}$. Less intuitive than relative risk, but in epidemiological studies we often don't know I and C (exposed/non-exposed).

Odds ratio The odds of an event A are defined as $P(A)/P(A^c)$; then $P(A) = \text{odds}/(1 + \text{odds})$. Trick: if odds are m/n , probability is $m/(m+n)$; if probability is m/n , odds are $m/(n-m)$.

The odds of $A|B$ are defined in the same way; the odds ratio of $A|B$ (= odds ratio of $B|A$) is defined as

$$\psi = \frac{\text{odds}(D|E)}{\text{odds}(D|E^c)} = \frac{P(D|E)/P(D^c|E)}{P(D|E^c)/P(D^c|E^c)} = \frac{p_{11}p_{00}}{p_{01}p_{10}}$$

It measures the association between the events: D and E are independent, positively or negatively associated iff ψ is equal to, greater than or smaller than 1 respectively.

It is estimated using one of 3 methods which give the same result:

- Multinomial sampling: we draw a sample from the population, then $X = (X_{ij})_{ij} \sim \mathcal{M}(n, p)$, use $\hat{p}_{ij} = X_{ij}/n$ hence

$$\hat{\psi} = X_{11}X_{00}/X_{01}X_{10}.$$

- Cohort sampling: we are given some exposed and non-exposed people and we count the number of people with the disease in both groups:

$$X_{01} \sim \text{binom}(x_{0\bullet}, P(D|E^c)),$$

$$X_{11} \sim \text{binom}(x_{1\bullet}, P(D|E))$$

so we can estimate $\hat{P}(D|E) = X_{11}/x_{1\bullet}$, $\hat{P}(D|E^c) = X_{01}/x_{0\bullet}$ which is enough to compute $\hat{\psi}$ which is the same as above.

- Case-control sampling: we are given some people with and without disease and we count how many were exposed.

$$X_{10} \sim \text{binom}(x_{\bullet 0}, P(E|D^c)),$$

$$X_{11} \sim \text{binom}(x_{\bullet 1}, P(E|D))$$

and, analogously as above, we find estimates of $P(E|D^c)$, $P(E|D)$; we compute the odds ratio of $E|D$ which is the same as for $D|E$ and yields the same $\hat{\psi}$ as before.

Note: $P(D|E) - P(D|E^c)$ can't be estimated in this setting because, using Bayes' theorem, we would need an estimate of $P(D)$ which we don't have. However if the disease is rare, the relative risk approximates the odds ratio (rare disease assumption):

$$\text{RR} = \frac{P(D|E)}{P(D|E^c)} \rightarrow \psi \text{ as } P(D) \rightarrow 0$$

In general, OR tends to exaggerate the association (is farther from 1) compared to RR; they are both 1 if there is no association.

Note: both cohort and case-control studies can be retrospective or prospective, depending on whether we look back at existing data (post hoc) or we enroll new participants during the study.

Causal inference If Y measures the health status of a patient, let Y^1, Y^0 equal Y respectively when the patient is treated or not (counterfactual model). Of course in real life we can only observe one of the two values for each patient. The average causal effect or average treatment effect

$$\theta = \text{ATE} = E(Y^1) - E(Y^0) = \frac{1}{n} \sum_{i=1}^n (Y_i^1 - Y_i^0)$$

where n is the total number of patients, can only be estimated.

If Y is binary we may alternatively use causal odds ratio $\frac{P(Y^1=1)/P(Y^1=0)}{P(Y^0=1)/P(Y^0=0)}$, causal relative risk $P(Y^1=1)/P(Y^0=1)$.

Define association as $\alpha = E(Y|X=1) - E(Y|X=0)$ (where $X=1, 0$ means receiving treatment or not, respectively). In general $\alpha \neq \theta$ (association is not causation); we may have $\alpha > 0$ with $\theta = 0$ or even $\theta < 0$. An association between X and Y can be created by an association between (Y^0, Y^1) and X , e.g. healthy user bias. However, if the subjects are assigned to treatment randomly, then $\alpha = \theta$ because X and (Y^0, Y^1) are independent; by the law of large numbers we have a consistent estimator

$$\begin{aligned} \hat{\theta} &= \hat{\alpha} = \hat{E}(Y|X=1) - \hat{E}(Y|X=0) \\ &= \frac{1}{n_T} \sum_{i=1}^n Y_i X_i - \frac{1}{n - n_T} \sum_{i=1}^n Y_i (1 - X_i) \end{aligned}$$

where $n_T = \sum_1^n X_i$ is the number of treated patients.

Suppose the treatment is not binary but $X \in \mathbb{R}$ represents the dose of a drug. We replace (Y^0, Y^1) with a function Y^x , the causal regression function is $\theta(x) = E(Y^x)$, the regression function is $r(x) = E(Y|X=x)$. Random assignment implies $\theta(x) = r(x)$; otherwise we may have constant θ (no causation) with non-constant r (association).

Adjusting for confounding If treatment is not assigned randomly, the study is called observational: the patients select their own value X of treatment. In this case we should control for confounding variables, i.e. find groupings of the subjects such that, within each group, subjects have similar characteristics and hence we may assume that X is essentially random, that is, X and Y^x are independent. Denote these groups with a random variable Z : then

$$\theta(x) = \int E(Y|X=x, Z=z) dF_Z(z)$$

(compare with $r(x) = E(Y|X=x) = \int E(Y|X=x, Z=z) dF_{Z|X}(z|x)$).

(A confounder is a variable causing the first and second variables. On the other hand, a mediator is caused by the first variable and causes the second.)

Simpson's paradox A treatment that is beneficial for men and beneficial for women seems to be harmful overall if we have $P(Y=1|X=1) < P(Y=1|X=0)$; the explanation is that this formula does not mean that the treatment is harmful, which should be written as $P(Y^1=1) < P(Y^0=1)$.

If in group i ($i=1, 2$) there are m_i admitted males over M_i male applicants, and analogously for females, then we may have $m_i/M_i < f_i/F_i$ in both groups

and still $\frac{m_1+m_2}{M_1+M_2} > \frac{f_1+f_2}{F_1+F_2}$ if more men apply to the group with highest admission rate. Analogously: let 1 and 2 be the groups of respectively young and old people, T_i , N_i denote respectively treated and non-treated status against a given disease, and t_i , n_i be the numbers of subjects in the two categories who have died from the disease: if treated people tend to be old ($N_1 \ll T_2, n_1 \ll t_2$), we may have $\frac{t_1+t_2}{T_1+T_2} > \frac{n_1+n_2}{N_1+N_2}$ even if the treatment works, and controlling for group will indeed reveal that it works.

2.7 Data reduction

We let $\mathbf{X} = (X_1, \dots, X_n)$ denote the random vector and $\mathbf{x} = (x_1, \dots, x_n)$ a particular sample.

A statistic $T(\mathbf{X})$ defines a form of data reduction or data summary of the variables X_1, \dots, X_n . The image of T defines a partition of the sample space according to the value of the statistic.

Sufficiency $T(\mathbf{X})$ is a sufficient statistic for θ if $P_\theta(\mathbf{X} = \mathbf{x} | T(\mathbf{X}) = t)$ does not depend on θ . (The data processing inequality is an equality in this case). Note: we assume $T(\mathbf{X})$ has a discrete distribution otherwise the condition has probability 0.

The sufficiency principle states that in this case, any inference about θ should depend on $T(\mathbf{X})$ and not on the particular sample: the statistic contains all the useful information about θ .

Indeed, the distribution on $T^{-1}(t)$ can be computed from the model without knowing θ , and we can use it to generate observations \mathbf{Y} with the same conditional distribution as \mathbf{X} . We want to show that they have the same unconditional distribution:

$$\begin{aligned} P_\theta(\mathbf{X} = \mathbf{x}) &= P_\theta(\mathbf{X} = \mathbf{x} \wedge T(\mathbf{X}) = T(\mathbf{x})) = \\ &= P(\mathbf{X} = \mathbf{x} | T(\mathbf{X}) = T(\mathbf{x})) P_\theta(T(\mathbf{X}) = T(\mathbf{x})) = \\ &= P(\mathbf{Y} = \mathbf{x} | T(\mathbf{X}) = T(\mathbf{x})) P_\theta(T(\mathbf{X}) = T(\mathbf{x})) = \\ &= P_\theta(\mathbf{Y} = \mathbf{x} \wedge T(\mathbf{X}) = T(\mathbf{x})) = P_\theta(\mathbf{Y} = \mathbf{x}) \end{aligned}$$

for all \mathbf{x}, θ .

To verify that a statistic is sufficient, we must only verify that $P_\theta(\mathbf{X} = \mathbf{x} | T(\mathbf{X}) = T(\mathbf{x}))$ does not depend on θ . By the equalities written above, this equals $p(\mathbf{x}|\theta)/q(T(\mathbf{x})|\theta)$, the ratio of the pmf's, so the condition is that this ratio is constant in θ .

A function $g(X_1, \dots, X_n, \theta)$ of a sample from a distribution depending on parameter θ is called a pivot if its distribution (but not necessarily the function itself) does not depend on θ . If it is also a statistic (i.e. if it can be computed just from the sample data, i.e. the whole function does not depend on unknown parameters), it is called an ancillary statistic. Such an unknown parameter is called nuisance parameter. Examples: if θ is the population mean, then the range $\max(X_1, \dots, X_n) - \min(X_1, \dots, X_n)$ and the sample variance are ancillary statistics, but not the sample mean, since its distribution depends on the population mean. Other example: Student's t-statistic for a Gaussian distribution with unknown mean and variance.

The ancillary complement of a statistic that is not sufficient is an ancillary statistic such that (T, U) is sufficient.

2.8 Bayesian statistics

Instead of minimizing the mean squared error of an estimator uniformly for θ , we minimize an average over θ . The posterior mean

$$\bar{\theta}_n = \int \theta f(\theta \mid x_1, \dots, x_n) d\theta$$

and a $1 - \alpha$ posterior confidence interval (a, b) by finding a and b such that the tails up to a and from b both have area $\alpha/2$, indeed

$$P(\theta \in (a, b) \mid x_1, \dots, x_n) = \int_a^b f(\theta \mid x_1, \dots, x_n) d\theta = 1 - \alpha.$$

Bayesian approach to curve fitting We introduce a prior distribution over the coefficients \mathbf{w} ; then we could find \mathbf{w} so as to maximize the posterior distribution $f(\mathbf{w} \mid \mathbf{x}, \mathbf{y}, \dots) \propto f(\mathbf{y} \mid \mathbf{x}, \mathbf{w}, \dots) f(\mathbf{w} \mid \dots)$, but this is equivalent to the old approach. Instead, in the Bayesian approach the predictive distribution is

$$f(y \mid x, \mathbf{x}, \mathbf{y}) = \int f(y \mid x, \mathbf{w}) f(\mathbf{w} \mid \mathbf{x}, \mathbf{y}) d\mathbf{w}$$

Chapter 3

Regression analysis

To study the relationship between X and Y we want to estimate the regression function, or conditional expectation

$$r(x) := E(Y|X = x) = \int yf(y|x)dy$$

using data $(Y_1, X_1), \dots, (Y_n, X_n) \sim F_{X,Y}$. Note that r is found by minimizing $E(Y - f(X))^2 = \int (y - f(x))^2 dF_{X,Y}(x, y)$ pointwise for each x .

The name *regression* was used by Galton and refers to the tendency of the response variable to "regress" towards the mean as the covariate approaches extreme values on either side.

3.1 Linear models

A linear model for n observations is, written in vector form,

$$\mathbf{y} = \mathbf{X}\beta + \varepsilon$$

(we may use a more general linear basis function model replacing \mathbf{X} with Φ). $\mathbf{X} \in \mathbb{R}^{n \times s}$ is the design matrix with full rank $s < n$ (this is required to invert $\mathbf{X}^t \mathbf{X}$, it means that there is no perfect multicollinearity i.e. dependence between the predictors; the rank is the number of degrees of freedom; if \mathbf{X} close to singular use singular value decomposition?), $\beta = (\beta_1, \dots, \beta_s)^t$ parameter vector to be estimated, $\varepsilon = (\varepsilon_1, \dots, \varepsilon_n)^t$ error vector with $\varepsilon_i \in \mathcal{L}^2$ with known distribution and mean zero, so

$$E(\mathbf{y}|\mathbf{X}) = \mathbf{X}\beta.$$

Assumptions The following assumptions are possible:

- Normal (or Gaussian) linear models: ε_i iid $\sim \mathcal{N}(0, \sigma^2)$.
- Gauss-Markov assumptions are more general:
 - zero-mean $E(\varepsilon_i) = 0$
 - homoscedastic, i.e. same finite variance $\text{Var}(\varepsilon_i) = v$

- non-autocorrelated $\text{Cov}(\varepsilon_i, \varepsilon_j) = 0$ if $i \neq j$, or equivalently $\text{E}(\varepsilon_i \varepsilon_j) = 0$ (though not necessarily independent).
- Generalized Gauss-Markov: $\text{E}(\varepsilon_i) = 0$, general covariance matrix (possibly correlated or heteroscedastic). Used in generalized least squares.

In a non-experimental science like econometrics we use instead the conditional means in our assumptions: $\text{E}(\varepsilon_i | \mathbf{x}_1, \dots, \mathbf{x}_n) = 0$ for $i = 1, \dots, n$ (exogeneity), which implies: $\text{E}(\varepsilon_i) = 0$ (by the law of total expectation), orthogonality $\text{E}(\mathbf{x}_j \cdot \varepsilon_i) = \mathbf{0}$, and $\text{Cov}(\varepsilon_i, x_{jk}) = 0$.

Note on multicollinearity Perfect multicollinearity is defined as a relationship $\lambda_0 + \sum_{i=1}^p \lambda_i X_i = 0$ where X_i is the vector of observations for the i -th parameter; general multicollinearity is defined adding an error term to the equation so that the relation can only hold approximately, which gives rise to invertible but possibly ill-defined matrices in the OLS model.

Signs of multicollinearity are big changes when modifying the predictors, insignificant coefficients for the affected variables despite positive joint F -tests or significant coefficients in the corresponding simple regression.

The statistical model is

$$(\mathbb{R}^n, \mathcal{B}^n, P_{\beta, v} \mid \beta \in \mathbb{R}^s, v > 0)$$

where $P_{\beta, v}$ is the distribution of the random vector $A\beta + \varepsilon$.

Different design matrices \mathbf{X} and error vectors give different models: a trivial example is the Gaussian product model $\mathcal{N}_{m, v}^{\otimes n}$ obtained when $\mathbf{X} = (1, \dots, 1)^t$, $\beta = m \in \mathbb{R}$, $\varepsilon \sim \mathcal{N}_n(0, \mathbb{I}v)$. Simple linear regression is given by $s = 2$,

$$\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_n \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \end{pmatrix} + \begin{pmatrix} \varepsilon_1 \\ \vdots \\ \varepsilon_n \end{pmatrix}$$

and analogously for multiple linear regression, polynomial regression. A Gaussian linear model is any linear model with $\varepsilon \sim \mathcal{N}_n(0, \mathbb{I}v)$.

(For multiple regression with k outputs to predict we replace \mathbf{y} with an $n \times k$ matrix \mathbf{Y} (each row being a transposed observation) and the vector β with a $s \times k$ matrix \mathbf{B}).

Ordinary least squares (OLS) An estimate $\hat{\beta}$ defines the fitted line $\hat{r}(X) = (1, X_1, \dots, X_p)\hat{\beta}$, the fitted values $\hat{y}_i = \hat{r}(\mathbf{x}_i)$, the residuals $\hat{\varepsilon}_i = y_i - \hat{y}_i$ and the residual sum of squares (also called sum of squared errors)

$$\text{RSS} = \sum_{i=1}^n \hat{\varepsilon}_i^2 = \|\mathbf{y} - \hat{\mathbf{y}}\|^2 = \|\mathbf{y} - \mathbf{X}\hat{\beta}\|^2$$

(or equivalently the $\text{MSE} = \text{RSS}/n$, or the $\text{RMSE} = \sqrt{\text{RSS}/n}$). The least squares estimate is the vector $\hat{\beta}$ minimizing RSS, i.e. the projection of \mathbf{y} onto the s -dimensional subspace $\text{Im}\mathbf{X} \subset \mathbb{R}^n$. This projection is $\Pi = \mathbf{X}(\mathbf{X}^t \mathbf{X})^{-1} \mathbf{X}^t$ and the solution is

$$\hat{\beta} := (\mathbf{X}^t \mathbf{X})^{-1} \mathbf{X}^t \mathbf{y}$$

found by solving: $\frac{\partial \text{RSS}}{\partial \beta_j} = 0$ for all j iff $-2^t(\mathbf{y} - \mathbf{X}\beta) = 0$ iff $\mathbf{X}^t\mathbf{X}\beta = \mathbf{X}^t\mathbf{y}$. Without a full statistical model, we only have a mathematical solution but not a statistical solution: we can't derive statistical properties (e.g. confidence intervals) of the estimator. So, we need to start adding assumptions.

Gauss-Markov assumptions (same finite unknown variance σ and uncorrelatedness of the ε_i). Then: $\hat{\beta}_i$ are consistent and the least-squares estimator is a BLUE (best linear unbiased estimator) of β , easily seen from the definition.

$$\mathbb{E}(\hat{\beta}) = \beta \quad \text{Var}(\hat{\beta} \mid \mathbf{x}_1, \dots, \mathbf{x}_n) = \sigma^2(\mathbf{X}^t\mathbf{X})^{-1}$$

Gaussian linear model If ε_i are normal, the least squares estimator is the MLE which can be found by maximizing the (log-)likelihood function.

$$f(\mathbf{y} \mid \beta, \sigma^2) = \prod_{i=1}^n \mathcal{N}_{\beta\mathbf{x}_i, \sigma^2}(y_i)$$

(note the use of independence).

If we estimate σ^2 , we can get the $\widehat{\text{SE}}(\hat{\beta}_i)$ from the covariance matrix. The sample variance RSS/n is a MLE for σ^2 , but in order to get an unbiased one we must correct it to

$$\hat{\sigma}^2 := \frac{\text{RSS}}{n-s} = \frac{\|\mathbf{y} - \mathbf{X}\hat{\beta}\|^2}{n-s} \quad (\text{for } \mathbf{X} \text{ as above } n-s = n-p-1)$$

indeed, we have $(n-s)V^* = \|\mathbf{y} - \mathbf{X}\hat{\beta}\|^2 = v\|\xi - \Pi\xi\|^2$. Using an orthonormal basis change of the argument vector, $\eta = O^t\xi$, this can be written $v\sum_{s+1}^n \eta_k^2$. Since the ξ_j are uncorrelated and standardized, we have $\mathbb{E}(\eta_k^2) = \mathbb{E}(\sum_{i,j} O_{ik}O_{jk}\xi_i\xi_j) = \sum_{i=1}^n O_{ik}^2 = 1$.

Tests for Gaussian linear models For the Gaussian model $\xi \sim N_n(0, \mathbb{1})$ so that $P_\theta = N_n(A\gamma, v\mathbb{1})$, so we can generalize Student's theorem 2.3:

$$\hat{\beta} \sim \mathcal{N}(\beta, \sigma^2(\mathbf{X}^t\mathbf{X})^{-1}) \quad \frac{n-s}{v}V^* \sim \chi_{n-s}^2$$

so we have the usual approximate $1 - \alpha$ confidence intervals for β_i using $\widehat{\text{SE}}(\hat{\beta}_i)$, and we have t-tests for $H_0: \beta_j = 0$ using the t -statistic

$$\frac{\hat{\beta}_j - \beta_j}{\hat{\sigma}\sqrt{v_j}} \sim t_{n-p-1} \quad (v_j = [(\mathbf{X}^t\mathbf{X})^{-1}]_{jj})$$

- a confidence region (ellipsoid) for β with error level α , and an approximate one, are respectively

$$C(-) = \{\beta \in \mathbb{R}^s : |A(\beta - \hat{\beta})|^2 < sV^*f_{s,n-s;1-\alpha}\} \quad (3.1)$$

$$C(-) = \{\beta \in \mathbb{R}^s : |A(\beta - \hat{\beta})|^2 < V^*\chi_{s;1-\alpha}^2\} \quad (3.2)$$

Indeed, $|\mathbf{X}(\hat{\beta} - \beta)|^2/v \sim \chi_s^2$ is independent of V^* , and therefore $|\mathbf{X}(\hat{\beta} - \beta)|^2/sV^* \sim F_{s,n-s}$.

- χ^2 -tests: an α -confidence interval for the variance is

$$C(-) = \left(\frac{n-s}{\chi_{n-s;1-\alpha/2}^2} V^*, \frac{n-s}{\chi_{n-s;\alpha/2}^2} V^* \right)$$

(follows immediately from the distribution of the rescaled variance above) and a test for $H_0: v \leq v_0$ (resp. \geq) has rejection region $(n-s)V^* > v_0 \chi_{n-s;1-\alpha}^2$ (resp. ?)

- t -tests: an α -confidence interval for a linear characteristic $\tau(\beta) = c \cdot \beta$ is

$$C(-) = (c \cdot \hat{\beta} \mp \delta \sqrt{V^*})$$

and a level- α test of $H_0: c \cdot \beta = m_0$ has rejection region $|c \cdot \hat{\beta} - m_0| > \delta \sqrt{V^*}$, where $\delta = t_{n-s;1-\alpha/2} \sqrt{c^t (A^t A)^{-1} c}$.

- F -test for $H_0: A\gamma \in H$, with $H \subset \text{Im } A =: L$ subspace of dimension $r < s$. Then $|\Pi_L X - \Pi_H X|^2 / v \sim \chi_{s-r}^2$ is independent of V^* , so the Fisher statistic

$$F_{H,L} := \frac{n-s}{s-r} \frac{|\Pi_L Y - \Pi_H Y|^2}{|Y - \Pi_L Y|^2} = \frac{|A\hat{\gamma} - \Pi_H Y|^2}{(s-r)V^*} \sim F_{s-r, n-s}$$

and a level- α test has rejection region $F_{H,L} > f_{s-r, n-s; 1-\alpha}$.

Generalized least squares (GLS) We are in the generalized Gauss-Markov assumptions, covariance matrix $\sigma^2 \Sigma$ where Σ is a known positive definite matrix and σ^2 is an unknown scale parameter; transform the data to $Y^* = \Sigma^{-1/2} Y$, $A^* = \Sigma^{-1/2} A$, then the new model falls into the ordinary case and we find its $\hat{\beta}$.

Model fit Too few covariates = high bias = underfitting. Too many covariates = high variance = overfitting.

R^2 coefficient of determination The most general definition is

$$R^2 = 1 - \frac{RSS}{TSS}.$$

For a linear regression model including a constant (i.e. design matrix has a column of 1's), we have

$$TSS = RSS + ESS$$

$$\sum (y_i - \bar{y})^2 = \sum (y_i - \hat{y}_i)^2 + \sum (\hat{y}_i - \bar{y})^2$$

Compare with ANOVA partition (the cross-term is zero by direct computation). $R^2 = \frac{ESS}{TSS}$ equals the squared sample correlation between the y_i and \hat{y}_i .

For simple linear regression we also have $ESS = S_{xy}^2 / S_{xx}$ where we defined $S_{xy} = \sum (x_i - \bar{x})(y_i - \bar{y})$, which provides the link to the t -test; $R^2 = \frac{S_{xy}^2}{S_{xx} S_{yy}}$ equals the squared sample correlation between the y_i and x_i .

Adjusted R^2 to select between models containing different numbers of variables, since the usual R^2 tends to decrease as we add more variables, we define the correction $1 - \frac{RSS/(n-d-1)}{TSS/(n-1)}$ (essentially we have replaced biased estimates of the variance, i.e. numerator and denominator both divided by n , with their unbiased corrections).

Assumptions We can check whether a linear model is appropriate by plotting the residuals and checking if they look like iid random variables. We should not discern patterns (e.g. a U-shape), heteroscedasticity (seen as an increase in magnitude of residuals; can be fixed by fitting first-order differences, log-transformed data or Box-Cox transformation) or autocorrelation (e.g. daily stock prices are autocorrelated, stock returns aren't).

Complements

- Regularization (weight decay/parameter shrinkage): one can control overfitting by limiting the size of the coefficients, if we replace the error function with $E_D(w) + \lambda E_W(w)$ i.e. we add a penalty term (governed by a to-be-chosen coefficient), usually of the form $\lambda \sum |w_j|^q$ for some $q > 0$.
- MLE might be difficult with big amounts of data, so we may want to use sequential learning (or on-line learning) techniques like stochastic descent to find the correct weights.
- Functional regression
- Error-in-variables (EIV) models: assume there are errors in the independent variables too, so X becomes a random variable like Y .
- Orthogonal least squares: minimizes orthogonal distances rather than vertical distances. This is not a MLE in general.

3.2 Analysis of variance (ANOVA)

Denotes methods to analyze the difference between means from different groups of observations, generalizing the t -test to more than two means. The most common form is based on linear models.

Observations, X_{ik} partitioned into groups from $G = \{1, \dots, s\}$ (one single family of groups = one-factor ANOVA); the index set is therefore $B = \{ik := (i, k) : i \in G, 1 \leq k \leq n_i\}$ with cardinality $n = \sum_{i \in G} n_i$.

We set up a model with an unknown mean vector $\beta = (m_i)_{i \in G}$ and an unknown parameter v such that $X_{ik} = m_i + \sqrt{v}\xi_{ik}$ which we write in matrix form as $X = A\beta + \sqrt{v}\xi$ where

$$A\beta = \begin{pmatrix} \mathbb{1}_{n_1} & & \\ & \ddots & \\ & & \mathbb{1}_{n_s} \end{pmatrix} \beta = (m_1, \dots, m_1, m_2, \dots, m_2, \dots, m_s, \dots, m_s)^T$$

The unbiased estimator $\hat{\beta}$ of $\beta = (m_i)_{i \in G}$ is just the vector of sample means:

$$\hat{\beta} = (A^T A)^{-1} (A^T X) = \begin{pmatrix} n_1 & & \\ & \ddots & \\ & & n_s \end{pmatrix}^{-1} \begin{pmatrix} n_1 M_1 & & \\ & \ddots & \\ & & n_s M_s \end{pmatrix} = \begin{pmatrix} M_1 \\ \vdots \\ M_s \end{pmatrix}$$

so that $A\hat{\beta} = (M_1, \dots, M_1, M_2, \dots, M_2, \dots, M_s, \dots, M_s)^T$ and

- $\hat{\sigma}^2 = \frac{\|X - A\hat{\beta}\|^2}{n-s} = \frac{1}{n-s} \sum_{ik \in B} (X_{ik} - M_i)^2 = \frac{1}{n-s} \sum_{i \in G} (n_i - 1) \hat{\sigma}_i^2 =: V_{\text{wg}}$,
the average sample variance within the groups;
- the total sample variance $V_{\text{tot}} = \frac{\|X - M\mathbb{1}_n\|^2}{n-1} = \frac{1}{n-1} \sum_{ik \in B} (X_{ik} - M)^2$;
- the variance between the groups

$$V_{\text{bg}} = \frac{\|A\hat{\beta} - M\mathbb{1}_n\|^2}{s-1} = \frac{1}{s-1} \sum_{i \in G} n_i (M_i - M)^2$$

is the empirical variance of the group means (with a weighting by the number of observations).

By Pythagoras' theorem we have the following (equivalent) relations

$$\begin{aligned} \|X - M\mathbb{1}_n\|^2 &= \|X - A\hat{\beta}\|^2 + \|A\hat{\beta} - M\mathbb{1}_n\|^2 \\ (n-1)V_{\text{tot}} &= (n-s)V_{\text{wg}} + (s-1)V_{\text{bg}} \end{aligned}$$

We also note that, unless the means of each group are all the same, V_{tot} is not an unbiased estimator of v :

$$\mathbb{E}_{\beta, v}(V_{\text{tot}}) = v + \frac{1}{n-1} \sum_{i \in G} n_i (m_i - \bar{m})$$

We assume ξ_{ik} are iid standard normal so we have the results obtained for Gaussian linear models: the confidence ellipsoid eq. (3.1) for $\beta = (m_i)_{i \in G}$, a t -test for the two-sample problem $m_1 = m_2$ is obtained for $s = 2$, $H_0: c \cdot \beta = 0$, $c = (1, -1)^T$, an F -test for a multi-sample problem $H_0: m_1 = \dots = m_s$ (we don't do $\binom{s}{2}$ tests because if α were chosen small enough to compensate, e.g. with a Bonferroni correction, the power would be too small) can be written as $H_0: A\beta \in H := \langle (1, \dots, 1)^T \rangle \subset \mathbb{R}^n$, $F_{H,L} = V_{\text{bg}}/V_{\text{wg}}$ so the rejection region is $V_{\text{bg}} > f_{s-1, n-s; 1-\alpha} V_{\text{wg}}$.

Chapter 4

Survival analysis

From a *failure density* $f(\tau)$ we define the failure distribution (its CDF, probability of failure before t) and the survival function (its complementary, the probability of survival up to t):

$$F(t) = P(T \leq t) = \int_0^t f(\tau) d\tau = 1 - S(t)$$

The instantaneous hazard rate (or failure rate), (also called force of mortality, instantaneous rate of mortality...), is

$$h(t) := \lim_{\Delta t \rightarrow 0} \frac{P(t \leq T \leq t + \Delta t \mid T \geq t)}{\Delta t} = \lim_{\Delta t \rightarrow 0} \frac{S(t) - S(t + \Delta t)}{\Delta t \cdot S(t)} = \frac{f(t)}{S(t)}.$$

Examples of common models

- Exponential density: $f(\tau) = \lambda e^{-\lambda\tau}$ has CDF $F(t) = 1 - e^{-\lambda t}$ and constant hazard rate $h(t) = \lambda$, showing the memorylessness of the exponential distribution.
- Gompertz' law of mortality: $h(t) = Bc^t$ with $B \in (0, 1)$ and $c > 1$.

Assumption If T_x denotes the future lifetime of a subject of age x , define F_x and S_x as above but with T_x instead of T . We use the assumption $P(T_x \leq t) = P(T_0 \leq x + t \mid T_0 > x)$, whence $S_0(x + t) = S_0(x)S_x(t)$. We have $h(x + t) = f_x(t)/S_x(t)$.

Actuarial notation

- ${}_t p_x := S_x(t) = P(T_x > t)$;
- ${}_t q_x := F_x(t) = P(T_x \leq t) = \int_0^t {}_s p_x \mu_{x+s} ds$, $\mu_x := h(x)$;
 $q_x := {}_1 q_x$ is called the mortality rate at age x .
- ${}_{u|t} q_x := P(u < T_x \leq u + t)$ (deferred mortality probability)
- The curtate future lifetime for a life aged x is the random variable $K_x := \lfloor T_x \rfloor$. By definition $P(K_x = k) = P(k \leq T_x < k + 1)$.

- ${}^{\circ}e_x := E(T_x)$ (complete expectation of life);
 $e_x := E(K_x) = \sum_{k=0}^{\infty} kP(K_x = k) = \sum_{k=1}^{\infty} S_x(k).$

Chapter 5

Design of experiments

Blocking, randomization, sampling, optimal experimental design, causal analysis, misuse/fallacies...

Chapter 6

Misc

Zipf-Mandelbrot law states that the relative frequency f of an item of rank k (the k th most frequent) in a corpus of N elements is often distributed according to the power-law distribution

$$f(k, N, q, s) = \frac{1/(k+q)^s}{H_{N,q,s}}$$

where $H_{N,q,s} = \sum_{i=1}^N (q+i)^s$ and $s > 0$ and q are parameters. For $q = 0$ and finite N this is Zipf's law. For $q = 0$ and $N = \infty$ this is a zeta distribution (if $s > 1$).

Benford's law is satisfied when, in a set of numbers, the leading digit (first digit) d is distributed according to

$$P(\{d\}) = \log_{10}(d+1) - \log_{10}(d).$$

This happens when the numbers are log-normally distributed, for example when they arise from multiplicative fluctuations, e.g. stock prices (whereas additive fluctuations give rise to normal distributions).